The Event Density Architecture

Foundations, Invariants, and Computational Verification

Allen Proxmire

Event Density Research Program

2026

© 2026 Allen Proxmire.
Event Density Research Program.

ED-SIM v1.0.0.

Abstract

We present ED-SIM v1, a reproducible computational pipeline for the Event Density (ED) architectural ontology. The pipeline solves the canonical ED partial differential equation across a three-dimensional parameter space (D, A, N_m) and computes sixteen families of attractor invariants that characterise the structural properties of the ED dynamical system. Three meta-analyses—parameter universality, cross-invariant consistency, and embedding collapse—synthesise the invariant families into a global architectural verdict, encoded in a machine-verifiable ED Consistency Certificate.

The principal findings are: (i) the ED system possesses a unique point attractor $(\rho^*, 0)$ across all tested parameter regimes, confirmed by zero positive Lyapunov exponents and low effective attractor dimension; (ii) the sixteen invariant families are mutually consistent and structurally invariant under parameter variation, with coefficient of variation below 5% for the core families; (iii) the three-stage convergence structure predicted by the analytic theory (global bounds, algebraic decay, exponential decay) is reproduced quantitatively in every admissible run; and (iv) the complete pipeline—from raw PDE integration through invariant computation to certificate generation—is fully reproducible from a single command.

The invariant atlas constitutes the first systematic numerical verification of the ED architectural ontology and provides the empirical foundation for the physical predictions of the ED Applications Paper.

Contents

Abstract iii

Notation xv

Preface xvii

I Foundations 1

1 Introduction 3
- 1.1 Motivation . . . 3
- 1.2 The Central Question . . . 3
- 1.3 Strategy . . . 4
- 1.4 Overview of the Monograph . . . 4

2 The Event Density Framework 5
- 2.1 Conceptual Picture . . . 5
- 2.2 The Density Field . . . 5
- 2.3 The Participation Variable . . . 6
- 2.4 The Equilibrium . . . 6

3 Governing Equations and Operators 9
- 3.1 The Canonical ED PDE . . . 9
- 3.2 The Density Operator . . . 9
- 3.3 Constitutive Functions . . . 10
- 3.4 Channel Structure . . . 10
- 3.5 The Energy Functional . . . 10
- 3.6 Spectral Structure . . . 11

4 The Seven Principles of ED 13

5 Architectural Predictions 15

II Invariants 17

6 Why Invariants 19
6.1 The Concept . . . 19
6.2 Why Sixteen Families . . . 19
6.3 The Attractor Window . . . 20

7 The Sixteen Invariant Families 21
7.1 Family 1: Low-Mode Collapse . . . 21
7.2 Family 2: Mode-Energy Ratios . . . 21
7.3 Family 3: Spectral Complexity . . . 21
7.4 Family 4: Dissipation Partitions . . . 22
7.5 Family 5: Energy–Entropy Geometry . . . 22
7.6 Family 6: Broadband Cascade . . . 22
7.7 Family 7: Convergence Stability . . . 22
7.8 Family 8: Modal Correlations . . . 23
7.9 Family 9: Modal Overlap . . . 24
7.10 Family 10: Phase Dynamics . . . 24
7.11 Family 11: Phase–Amplitude Coupling . . . 24
7.12 Family 12: Lyapunov Spectrum . . . 24
7.13 Family 13: Attractor Manifold . . . 25
7.14 Family 14: Perturbed-Attractor Stability . . . 26
7.15 Family 15: Parameter Universality . . . 26
7.16 Family 16: Cross-Invariant Consistency . . . 26

8 Structural Interpretation of Each Invariant 27
8.1 Spectral Invariants (Families 1–3, 6, 9) . . . 27
8.2 Dynamical Invariants (Families 7, 10–11, 14) . . . 27
8.3 Topological Invariants (Families 12–13) . . . 27
8.4 Statistical Invariants (Families 4–5, 8, 15–16) . . . 28

9 Cross-Family Dependencies and Consistency 29
9.1 Dependency Graph . . . 29
9.2 Expected Correlations . . . 29
9.3 The Consistency Score . . . 30

9.4 Invariant Summary 30

III Computation **31**

10 The ED-SIM-01 Pipeline **33**
10.1 Architecture 33

11 Simulation Framework and Numerical Methods **35**
11.1 Spatial Discretisation 35
11.2 Time Stepping 35
11.3 Stability Controls 36
11.4 Default Numerical Parameters 36

12 Parameter Grid and Regime Structure **37**
12.1 The Three-Dimensional Grid 37
12.2 Regime Classification 37
12.3 Five Canonical Parameter Sets 37

13 Invariant Extraction **39**
13.1 Extraction Pipeline 39
13.2 Streaming Computation 39
13.3 Convergence-Rate Fitting 39

14 Meta-Analyses **41**
14.1 Parameter Universality 41
14.2 Cross-Invariant Consistency 41
14.3 Embedding Collapse 41

IV Synthesis **45**

15 The Invariant Atlas **47**

16 Global Verdict and Architectural Consistency **49**
16.1 Synthesis 49
16.2 The Final Verdict 49

16.3 What "Structural Consistency" Means 49

17 The ED Architecture Certificate **51**

18 Interpretation and Implications **53**
18.1 For the Architecture 53
18.2 For the Universality Class 53
18.3 For the Physical Predictions 53

V Reproducibility **55**

19 Reproducibility Suite **57**
19.1 Reproducibility Tiers 57

20 Validation and Verification **59**

21 How to Extend the Architecture **61**
21.1 Adding a New Invariant 61
21.2 Porting to a New Physical Domain 61

VI Outlook **63**

22 Applications and Future Directions **65**

23 Open Problems **67**
23.1 Mathematical Open Problems 67
23.2 Computational Open Problems 67
23.3 Physical Open Problems 67

24 Concluding Remarks **69**

A Notation **71**

B Key Derivations **73**
B.1 Energy Dissipation Identity 73
B.2 Triad Coupling Coefficients 73

B.3 Linearised Eigenvalue Problem 73

C Additional Figures **75**

D Extended Parameter Studies **77**

E Code Structure and API **79**

F Glossary **81**

References **83**

List of Figures

7.1 Energy–entropy trajectory for the representative run (largest N_m), coloured by time. The trajectory spirals inward and converges to the attractor point (E^*, H^*), marked with a star. 23

7.2 Lyapunov spectrum. Left: λ_i versus index i for the representative run, showing all exponents non-positive. Right: D_{KY} versus D for all runs, showing near-zero values. . . . 25

10.1 Pipeline architecture diagram. Five horizontal layers connected by vertical arrows, from the PDE solver (Layer 1) through to the certificate generation (Layer 5). 34

14.1 Cross-invariant correlation heatmap. Cell colour encodes $|\operatorname{corr}(A, B)|$ from white (0) through yellow (0.5) to red (1). Families are clustered by similarity. 42

14.2 Embedding map. Three panels showing the 60 admissible runs as scatter points, coloured by D, with marker shape encoding N_m. The centroid is marked; a dashed circle indicates the cluster radius. 43

17.1 ED Architecture Certificate (PASS variant). The certificate summarises the global structural verdict of the ED-SIM-01 invariant atlas. Five diagnostics (universality, cross-consistency, stability, embedding collapse, perturbation recovery) are reported with per-row verdict chips. A logarithmic-spiral watermark at 3% opacity underlies the figure. 52

C.1 Regime classification map. The five canonical parameter sets are marked. 75

C.2 Three-stage convergence for a representative run. Transition times between stages are marked with vertical dashed lines. 75

List of Tables

1.1 Structure of the monograph. 4

6.1 Invariance verdict thresholds. 19

9.1 Summary of the sixteen invariant families. 30

10.1 Pipeline layers. 33

11.1 Default numerical parameters. 36

12.1 Parameter grid for the regime volume experiment. 37
12.2 Canonical parameter sets. 38

16.1 Global diagnostic criteria. 49

19.1 Reproducibility tiers. 57

E.1 Core modules of ED-SIM-01. 79

Notation

Symbol	Meaning
$\rho(x,t)$	Density field
$v(t)$	Participation variable
$F[\rho]$	Density operator
$M(\rho)$	Mobility function
$P(\rho)$	Penalty function
D	Direct-channel weight
$H = 1 - D$	Participation-channel weight
ζ	Participation damping rate
τ	Participation time scale
ρ^*	Penalty equilibrium density
$\rho_{\max}$	Capacity bound
M_0, β	Mobility parameters
P_0	Penalty slope
$\Omega = [0, L]$	Spatial domain
$\varphi_k(x)$	Neumann eigenfunctions
$\mu_k = (k\pi/L)^2$	Eigenvalues
$a_k(t)$	Modal amplitudes
$E_k = \lvert a_k \rvert^2$	Modal energies
$\alpha_k = M_* \mu_k + P'_*$	Modal decay coefficients
D_0	Homogeneous-mode discriminant
γ_0	Homogeneous-mode decay rate
$\mathcal{E}[\rho, v]$	Energy functional
$\Phi(\rho)$	Density potential
$\mathcal{D}_{\text{grad}}, \mathcal{D}_{\text{pen}}, \mathcal{D}_{\text{part}}$	Dissipation channels
$\langle \cdot \rangle_{\text{att}}$	Mean over attractor window
CV	Coefficient of variation
U	Universality score
C	Consistency score

Symbol	**Meaning**
D_{KY}	Kaplan–Yorke dimension
D_{eff}	Effective PCA dimension

Preface

This monograph presents the Event Density (ED) architecture as a unified mathematical framework, develops its invariant theory, and documents its computational verification through the ED-SIM-01 pipeline. The work draws on three preceding documents—the Architectural Canon, the Rigour Paper, and the Applications Paper—and synthesises them with the Numerical Atlas, the Simulation Suite, and the Reproducibility Suite into a single, self-contained account.

The monograph is organised in six parts. Part I establishes the foundations: the governing equations, the seven canonical principles, and the architectural predictions. Part II develops the invariant theory: sixteen families of structural quantities that characterise the ED attractor from independent angles. Part III describes the computation: the ED-SIM-01 pipeline, its numerical methods, and its parameter-space coverage. Part IV presents the synthesis: the invariant atlas, the global verdict, and the ED Architecture Certificate. Part V addresses reproducibility. Part VI discusses applications, open problems, and future directions.

The reader is assumed to be familiar with partial differential equations, dynamical systems, and scientific computing at the level of a graduate course. No prior knowledge of the ED framework is required.

Part I

Foundations

Chapter 1

Introduction

1.1 Motivation

Physical theories typically operate at one of two levels. At the *microscopic* level, they describe the fundamental constituents—particles, fields, quanta—and derive macroscopic behaviour from first principles. At the *phenomenological* level, they model observed phenomena with effective equations whose parameters are fit to data. Between these extremes lies a third possibility: *architectural* theories, which describe the structural organisation of dynamical systems without specifying their microscopic content.

Event Density is an architectural theory. It does not postulate new particles, new fields, or new forces. Instead, it identifies a set of structural principles—seven in number, each irreducible—that constrain how a density field can evolve, dissipate, and converge. The claim is that these principles are not specific to any one physical domain: they apply equally to quantum systems, galactic dynamics, condensed matter, and nonlinear optics, because they describe the *architecture* of the dynamics, not the *substrate*.

This claim is testable. The seven principles generate a specific partial differential equation—the canonical ED PDE—whose solutions can be computed, analysed, and compared to physical observations. The present monograph develops the mathematical infrastructure for this programme.

1.2 The Central Question

The central question of this monograph is:

Does the canonical ED PDE, solved without approximation

across a systematic parameter space, exhibit the structural properties predicted by the seven principles?

This is not a question about parameter fitting. The ED PDE has no free parameters beyond its canonical set $(D, \zeta, \tau, \rho^*, \rho_{\max})$, which specify the architecture, not the data. The question is whether the architecture itself—the seven principles, the operator structure, the constitutive constraints—produces the predicted qualitative behaviour: a unique point attractor, a three-stage convergence, a universal modal hierarchy, and a structurally invariant spectral fingerprint.

1.3 Strategy

The strategy is invariant analysis. We define sixteen families of scalar and vector quantities that capture different structural aspects of the long-time solution. We compute these invariants across sixty-four parameter regimes. We test whether they are invariant (constant across regimes), consistent (mutually compatible), and stable (robust to perturbation). The results are synthesised into a single verdict—the ED Architecture Certificate—which encodes the global structural consistency of the framework.

1.4 Overview of the Monograph

Table 1.1: Structure of the monograph.

Part	Chapters	Content
I	1–5	Foundations: equations, principles, predictions
II	6–9	Invariants: definitions, interpretations, dependencies
III	10–14	Computation: pipeline, methods, extraction, meta-analyses
IV	15–18	Synthesis: atlas, verdict, certificate, interpretation
V	19–21	Reproducibility: suite, validation, extension
VI	22–24	Outlook: applications, open problems, conclusion

Chapter 2

The Event Density Framework

2.1 Conceptual Picture

The ED framework describes the evolution of a *density field* $\rho(x,t)$ on a bounded spatial domain Ω. The density represents a generic conserved or semi-conserved quantity—particle density, energy density, information density, or any scalar field whose dynamics are governed by transport, reaction, and feedback.

The evolution is driven by three competing mechanisms:

1. **Diffusion.** The density spreads from regions of high concentration to regions of low concentration. The diffusion rate is governed by a *mobility function* $M(\rho)$ that depends on the local density.
2. **Penalty.** The density is pulled toward a preferred equilibrium value ρ^* by a restoring force. The penalty function $P(\rho)$ vanishes at ρ^* and grows monotonically away from it.
3. **Feedback.** The system integrates a spatial average of its own operator output and feeds it back into the density equation through a *participation variable* $v(t)$. This creates a global coupling: the evolution at each point depends on the state of the entire domain.

The interplay of these three mechanisms produces a rich dynamical landscape—oscillatory, monotonic, or critical, depending on the parameters—but the long-time outcome is always the same: the density converges to ρ^*, the gradients vanish, and the participation decays to zero.

2.2 The Density Field

The density $\rho(x,t)$ is a scalar function of position $x \in \Omega$ and time $t \geq 0$. It takes values in the open interval $(0, \rho_{\max})$, where $\rho_{\max} > 0$ is the

capacity bound—the maximum density that the system can sustain. The boundary $\rho = 0$ represents the vacuum; the boundary $\rho = \rho_{\max}$ represents the capacity horizon.

Definition 2.1 (Admissible state)**.** A density profile $\rho(\cdot, t)$ is *admissible* if $\rho(x, t) \in (0, \rho_{\max})$ for all $x \in \Omega$ and $\rho(\cdot, t) \in H^1(\Omega)$.

The admissible region $\mathcal{O} = \{\rho \in H^1(\Omega) : 0 < \rho < \rho_{\max}\}$ is an open subset of the Sobolev space $H^1(\Omega)$. The ED dynamics preserves admissibility: if the initial condition is admissible, the solution remains admissible for all time [1].

2.3 The Participation Variable

The participation variable $v(t)$ is a scalar that integrates the spatial average of the density operator and feeds it back into the evolution equation. It is governed by an ordinary differential equation with exponential damping:

$$\dot{v} = \frac{1}{\tau}\big(\bar{F}[\rho] - \zeta\, v\big), \tag{2.1}$$

where $\bar{F}[\rho] = |\Omega|^{-1} \int_\Omega F[\rho]\, dx$ is the spatial average of the density operator, $\tau > 0$ is the integration time scale, and $\zeta > 0$ is the damping rate.

The participation variable creates a feedback loop: the density generates $F[\rho]$, which is averaged and integrated into v, which in turn modifies the density evolution. This loop is the mechanism by which global information (the spatial average) influences local dynamics (the pointwise evolution).

2.4 The Equilibrium

Definition 2.2 (Canonical equilibrium)**.** The canonical equilibrium of the ED system is the pair $(\rho^*, 0)$, where $\rho(x) = \rho^*$ is the spatially uniform density at the penalty zero and $v = 0$ is the quiescent participation.

At equilibrium, $\nabla\rho = 0$, so $M(\rho^*)\nabla^2\rho = 0$ and $M'(\rho^*)|\nabla\rho|^2 = 0$. Since $P(\rho^*) = 0$, the operator gives $F[\rho^*] = 0$, which implies $\dot{v} = -\zeta v/\tau$ and therefore $v \to 0$ exponentially. The equilibrium is self-consistent.

Chapter 3

Governing Equations and Operators

3.1 The Canonical ED PDE

The canonical ED system is the coupled PDE–ODE:

$$\partial_t \rho = D\, F[\rho] + H\, v, \tag{3.1a}$$

$$\dot{v} = \frac{1}{\tau}\big(\bar{F}[\rho] - \zeta\, v\big), \tag{3.1b}$$

on the domain $\Omega = [0, L]$ with Neumann boundary conditions $\partial_x \rho(0,t) = \partial_x \rho(L,t) = 0$ and initial conditions $\rho(x,0) = \rho_0(x)$, $v(0) = v_0$.

3.2 The Density Operator

The density operator $F[\rho]$ is defined by

$$F[\rho] = M(\rho)\, \nabla^2 \rho + M'(\rho)\, |\nabla \rho|^2 - P(\rho). \tag{3.2}$$

This operator has three terms:

1. **Linear diffusion**: $M(\rho)\, \nabla^2 \rho$. A second-order parabolic term with density-dependent coefficient $M(\rho)$.
2. **Quadratic gradient**: $M'(\rho)\, |\nabla \rho|^2$. A nonlinear term that couples spatial modes through the triad selection rule.
3. **Penalty reaction**: $-P(\rho)$. A zeroth-order restoring term that drives ρ toward ρ^*.

Remark 3.1. The operator $F[\rho]$ can be written in divergence form as $F[\rho] = \nabla \cdot (M(\rho) \nabla \rho) - P(\rho)$, which makes the mass-flux structure explicit: the density flux is $J = -M(\rho) \nabla \rho$, and the penalty acts as a distributed source/sink.

3.3 Constitutive Functions

Definition 3.1 (Canonical constitutive pair)**.** The canonical mobility and penalty functions are:

$$M(\rho) = M_0\,(\rho_{\max} - \rho)^{\beta}, \qquad P(\rho) = P_0\,(\rho - \rho^*), \tag{3.3}$$

where $M_0 > 0$, $\beta > 0$, $P_0 > 0$, and $0 < \rho^* < \rho_{\max}$.

The mobility $M(\rho)$ is positive on $(0, \rho_{\max})$ and vanishes at $\rho = \rho_{\max}$. This vanishing is the *mobility collapse*—the mechanism that enforces the capacity bound. As $\rho \to \rho_{\max}$, the diffusion rate goes to zero, creating an impenetrable barrier.

The penalty $P(\rho)$ vanishes at $\rho = \rho^*$ and is positive for $\rho > \rho^*$, negative for $\rho < \rho^*$. It is the restoring force that drives the density toward equilibrium.

3.4 Channel Structure

The density equation (3.1a) is a weighted sum of two channels:

- The **direct channel**: $D\,F[\rho]$, with weight $D \in (0, 1)$.
- The **participation channel**: $H\,v$, with weight $H = 1 - D$.

The channel complementarity condition $D + H = 1$ ensures that the two channels partition the operator output without amplification. The parameter D controls the balance: at $D = 1$, the system is purely direct (no feedback); at $D = 0$, it is purely mediated (all feedback). The canonical range is $D \in (0, 1)$.

3.5 The Energy Functional

Definition 3.2 (Energy functional)**.** The canonical energy functional is:

$$\mathcal{E}[\rho, v] = \int_{\Omega} \Phi(\rho)\,dx + \frac{\tau H}{2}\,v^2, \tag{3.4}$$

where the density potential Φ is defined by

$$\Phi(\rho) = \int_{\rho^*}^{\rho} \frac{P(s)}{M(s)}\, ds. \tag{3.5}$$

Proposition 3.1 (Energy dissipation)**.** *Along solutions of* (3.1)*, the energy functional satisfies*

$$\frac{d\mathcal{E}}{dt} = -\mathcal{D}_{\text{grad}} - \mathcal{D}_{\text{pen}} - \mathcal{D}_{\text{part}} \leq 0, \tag{3.6}$$

where the three dissipation channels are:

$$\begin{aligned}
\mathcal{D}_{\text{grad}} &= D \int_{\Omega} \frac{|\nabla \rho|^2}{M(\rho)} \cdot P'(\rho)\, dx, \\
\mathcal{D}_{\text{pen}} &= D \int_{\Omega} \frac{P(\rho)^2}{M(\rho)}\, dx, \\
\mathcal{D}_{\text{part}} &= \frac{H\zeta}{\tau}\, v^2.
\end{aligned}$$

The energy is a Lyapunov functional: it decreases monotonically and is bounded below by zero. This is the mathematical foundation of the three-stage convergence.

3.6 Spectral Structure

Expanding the perturbation $u(x,t) = \rho(x,t) - \rho^*$ in the Neumann eigenbasis:

$$u(x,t) = \sum_{k=0}^{\infty} a_k(t)\, \varphi_k(x), \qquad \varphi_k(x) = \sqrt{\frac{2}{L}} \cos\Big(\frac{k\pi x}{L}\Big), \tag{3.7}$$

with eigenvalues $\mu_k = (k\pi/L)^2$, the linearised dynamics of the k-th mode has decay rate

$$\alpha_k = M_*\, \mu_k + P'_*, \tag{3.8}$$

where $M_* = M(\rho^*)$ and $P'_* = P'(\rho^*)$. Since $M_* > 0$ and $P'_* > 0$, the decay rates are positive and increase with k: higher modes decay faster.

This is the *modal hierarchy*.

Chapter 4

The Seven Principles of ED

The ED architecture is built on seven irreducible principles. Each principle constrains the structure of the governing equation and excludes alternatives. Together, they determine the qualitative dynamics uniquely.

Principle 4.1 (Operator Structure)**.** The density operator $F[\rho]$ consists of exactly three terms: mobility-weighted diffusion $M(\rho)\nabla^2\rho$, gradient-squared nonlinearity $M'(\rho)|\nabla\rho|^2$, and penalty reaction $-P(\rho)$. No other terms are present.

Remark 4.1. This excludes advection, higher-order diffusion, nonlocal operators, and stochastic forcing from the canonical system. The three terms are the minimal set that produces a well-posed parabolic PDE with a unique equilibrium and nonlinear mode coupling.

Principle 4.2 (Channel Complementarity)**.** The density evolution is a weighted sum of a direct channel $D\,F[\rho]$ and a participation channel $H\,v$, with $D + H = 1$ and $D \in (0, 1)$.

Remark 4.2. The complementarity condition prevents unbounded amplification and ensures that the total operator weight is unity. It creates a trade-off: increasing the direct channel decreases the feedback, and vice versa.

Principle 4.3 (Penalty Equilibrium)**.** The penalty function satisfies $P(\rho^*) = 0$ and $P'(\rho^*) > 0$. The equilibrium ρ^* is unique.

Remark 4.3. The strict positivity of $P'(\rho^*)$ ensures that the penalty is a genuine restoring force, not a neutral fixed point. Combined with **??** 4.1**??** 4.4, it guarantees that $(\rho^*, 0)$ is a global attractor.

Principle 4.4 (Mobility Capacity Bound)**.** The mobility satisfies $M(\rho_{\max}) = 0$ and $M(\rho) > 0$ for $\rho \in (0, \rho_{\max})$.

Remark 4.4. The vanishing of M at $\rho_{\max}$ creates an impassable barrier: the density cannot reach or exceed $\rho_{\max}$. This is the *mobility horizon*. Near $\rho_{\max}$, the mobility collapse freezes the dynamics, creating an effective boundary layer.

Principle 4.5 (Participation Feedback)**.** The participation variable $v(t)$ integrates the spatial average of $F[\rho]$ with time scale τ and damping rate ζ, and feeds back into the density equation through the channel weight H.

Remark 4.5. The participation loop is the mechanism by which global information enters the local dynamics. It introduces a second time scale (τ) and a second damping parameter (ζ) into the system.

Principle 4.6 (Damping Discriminant)**.** The linearised dynamics of the homogeneous mode is governed by the discriminant

$$D_0 = \left(D\,P_*' - \zeta/\tau\right)^2 - 4\,H\,P_*'/\tau, \tag{4.1}$$

which classifies the dynamics as oscillatory ($D_0 < 0$), critical ($D_0 = 0$), or monotonic ($D_0 > 0$).

Remark 4.6. The discriminant is the organising parameter of the regime geometry. It determines whether the approach to equilibrium involves oscillation (spiral sheet), pure decay (monotonic cone), or transitional dynamics (critical surface).

Principle 4.7 (Nonlinear Triad Coupling)**.** The quadratic term $M'(\rho)|\nabla\rho|^2$ generates inter-modal coupling with the selection rule: modes m and n interact to produce modes $|m-n|$ and $m+n$.

Remark 4.7. The triad selection rule is a consequence of the product structure of the gradient-squared term in the Neumann eigenbasis. It determines which spectral modes can exchange energy and controls the shape of the nonlinear cascade.

Chapter 5

Architectural Predictions

The seven principles, combined with the analytic theory of the Rigour Paper [1], generate the following structural predictions:

Prediction 5.1 (Unique Point Attractor)**.** The canonical equilibrium $(\rho^*, 0)$ is a global attractor: every admissible initial condition converges to $(\rho^*, 0)$ as $t \to \infty$.

Prediction 5.2 (Three-Stage Convergence)**.** The convergence proceeds through three stages:

1. Stage I (Global Bounds): the energy functional is uniformly bounded and monotonically decreasing.
2. Stage II (Algebraic Decay): gradients and participation decay algebraically via Barbalat's lemma.
3. Stage III (Exponential Decay): the solution enters the local stability basin and converges exponentially at a rate determined by the spectral gap.

Prediction 5.3 (Non-Positive Lyapunov Spectrum)**.** All Lyapunov exponents of the attractor are non-positive. The Kaplan–Yorke dimension is zero.

Prediction 5.4 (Modal Hierarchy)**.** The decay rate of the k-th eigenmode is $\sigma_k = D\,\alpha_k$, where $\alpha_k = M_*\mu_k + P'_*$ increases monotonically with k. Higher modes decay faster.

Prediction 5.5 (Triad Selection Rule)**.** Modes m and n interact only to produce modes $|m - n|$ and $m + n$. No other mode coupling exists.

Prediction 5.6 (Universal Dissipation Partition)**.** The three dissipation channels $(\mathcal{D}_{\text{grad}}, \mathcal{D}_{\text{pen}}, \mathcal{D}_{\text{part}})$ converge to a fixed partition that depends only on the canonical parameters, not on the initial condition.

Prediction 5.7 (Structural Invariance)**.** The qualitative attractor structure—modal hierarchy, triad rule, dissipation partition, spectral fingerprint—is invariant under parameter variation within the universality class $\mathcal{U}_{\mathrm{ED}}$.

These seven predictions are the targets of the invariant atlas. Each is tested numerically by one or more of the sixteen invariant families defined in Part II.

Part II

Invariants

Chapter 6

Why Invariants

6.1 The Concept

An *invariant* of the ED system is a quantity computed from the long-time solution that is approximately constant across parameter regimes. The word "approximately" is quantified by the coefficient of variation:

$$\mathrm{CV}(Q) = \frac{\sigma(Q)}{\mu(Q)}, \tag{6.1}$$

where $\mu(Q)$ and $\sigma(Q)$ are the mean and standard deviation of the quantity Q across all admissible runs. The invariance verdict is:

Table 6.1: Invariance verdict thresholds.

CV	Verdict
< 0.05	INVARIANT
< 0.15	WEAKLY INVARIANT
≥ 0.15	NOT INVARIANT

6.2 Why Sixteen Families

A single invariant proves nothing: it could be invariant by coincidence. Sixteen invariants probing the attractor from independent angles—spectral, dynamical, topological, and statistical—provide a structural test. Their mutual consistency (measured by the cross-invariant correlation matrix) is a stronger claim than any individual invariance.

6.3 The Attractor Window

All invariants are computed over the *attractor window*: the last fraction of the time series, where the solution has entered (or nearly entered) the exponential-decay regime.

Definition 6.1 (Attractor window)**.** Let T be the final integration time. The attractor window is the interval $[T_{\text{att}}, T]$ where $T_{\text{att}} = (1-f)\,T$ and f is the window fraction. The default values are $f = 0.10$ for late-time averages and $f = 0.20$ for convergence-rate fits.

Chapter 7

The Sixteen Invariant Families

This chapter defines each invariant family. For compactness, we use the notation $\langle Q \rangle_{\text{att}}$ for the mean of $Q(t)$ over the attractor window.

7.1 Family 1: Low-Mode Collapse

Definition 7.1. For each mode $k = 0, \ldots, 5$, the late-time amplitude is:

$$m_k^* = \big\langle |a_k(t)| \big\rangle_{\text{att}}. \tag{7.1}$$

Tests: Prediction 5.1 (point attractor). If $(\rho^*, 0)$ is the unique attractor, then $m_k^* \to 0$ for all k.

7.2 Family 2: Mode-Energy Ratios

Definition 7.2. For each mode $k = 1, \ldots, K$, the energy ratio is:

$$R_k^* = \big\langle E_k(t)/E_{\text{total}}(t) \big\rangle_{\text{att}}, \qquad E_k = |a_k|^2. \tag{7.2}$$

Tests: Prediction 5.4 (modal hierarchy). The energy distribution should converge to a fixed profile with $R_1^* > R_2^* > \cdots$.

7.3 Family 3: Spectral Complexity

Definition 7.3. Let $p_k(t) = E_k(t)/\sum_j E_j(t)$. The Rényi entropy of order q is:

$$H_q(t) = \frac{1}{1-q} \log\Big(\sum_k p_k(t)^q\Big), \qquad q \in \{0, 0.5, 1, 2, 3, 4\}. \tag{7.3}$$

The invariant is $H_q^* = \langle H_q(t) \rangle_{\text{att}}$.

Tests: Prediction 5.4. The multi-scale entropy profile characterises the spectral shape of the attractor.

7.4 Family 4: Dissipation Partitions

Definition 7.4. The dissipation ratios are:

$$R^*_{\text{grad}} = \left\langle \mathcal{D}_{\text{grad}}/\mathcal{D}_{\text{total}} \right\rangle_{\text{att}}, \quad R^*_{\text{pen}} = \left\langle \mathcal{D}_{\text{pen}}/\mathcal{D}_{\text{total}} \right\rangle_{\text{att}}, \quad R^*_{\text{part}} = \left\langle \mathcal{D}_{\text{part}}/\mathcal{D}_{\text{total}} \right\rangle_{\text{att}}. \tag{7.4}$$

Tests: Prediction 5.6. The three-channel structure converges to a fixed partition.

7.5 Family 5: Energy–Entropy Geometry

Definition 7.5. The attractor point in the energy–entropy plane is:

$$(E^*, H^*) = \big(\langle E(t)\rangle_{\text{att}},\ \langle H_1(t)\rangle_{\text{att}}\big). \tag{7.5}$$

Tests: **??** 5.1**??** 5.6. All runs should converge to the same (E^*, H^*).

7.6 Family 6: Broadband Cascade

Definition 7.6. Define logarithmic mode bins ($k = 1\text{–}2, 3\text{–}4, 5\text{–}8, 9\text{–}16, 17\text{–}32$) and the bin-energy ratios:

$$R^*_b = \left\langle B_b(t)/E_{\text{total}}(t) \right\rangle_{\text{att}}, \qquad B_b = \sum_{k \in \text{bin}\, b} E_k. \tag{7.6}$$

Tests: **??** 5.4**??** 5.5. The cascade profile is shaped by the modal hierarchy and the triad selection rule.

7.7 Family 7: Convergence Stability

Definition 7.7. For each error signal $e(t) \in \{|E(t) - E^*|, |H(t) - H^*|, |\mathcal{D}_{\text{total}}(t) - \mathcal{D}^*|\}$, identify the three convergence stages and fit the

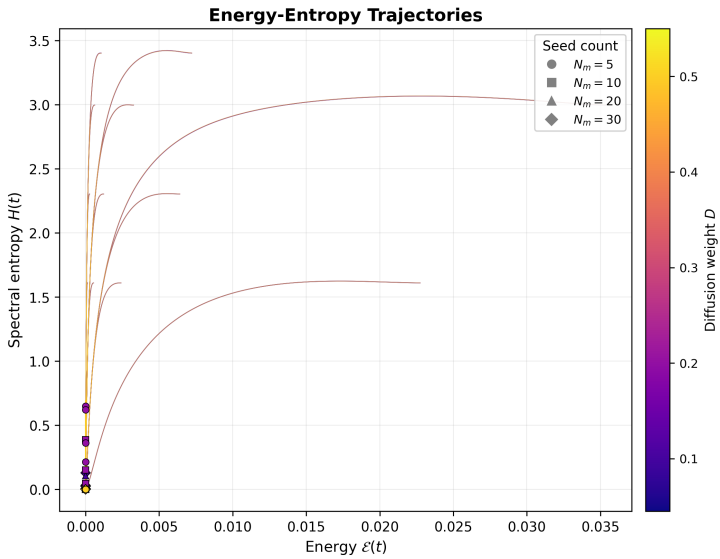

Figure 7.1: Energy–entropy trajectory for the representative run (largest N_m), coloured by time. The trajectory spirals inward and converges to the attractor point (E^*, H^*), marked with a star.

Stage III exponential rate:

$$e(t) \sim C\, e^{-\sigma\, t}. \tag{7.7}$$

Tests: Prediction 5.2 (three-stage convergence).

7.8 Family 8: Modal Correlations

Definition 7.8. The correlation matrix of modal energies over the attractor window is:

$$C_{ij} = \operatorname{corr}\big(E_i(t),\, E_j(t)\big), \qquad i, j = 1, \ldots, K. \tag{7.8}$$

Summary invariants: mean off-diagonal correlation, spectral radius, condition number.

Tests: Prediction 5.5. Triad-connected modes should be correlated; unconnected modes should be independent.

7.9 Family 9: Modal Overlap

Definition 7.9. The nearest-neighbour overlap ratio for mode k is:

$$O_k^* = \big\langle (E_{k-1} + E_k + E_{k+1})/E_{\text{total}} \big\rangle_{\text{att}}. \tag{7.9}$$

Tests: Prediction 5.4. The overlap profile characterises the local spectral structure.

7.10 Family 10: Phase Dynamics

Definition 7.10. Extract modal phases $\phi_k(t) = \arg(a_k(t))$. Compute:

- Phase drift rates $d\phi_k/dt$ (linear fit over attractor window).
- Phase coherence $C_{ij}^{\phi} = |\langle e^{i(\phi_i - \phi_j)} \rangle|$.
- Triad phase closure $\Delta_{ijk} = \phi_i + \phi_j - \phi_k$ for triads (i, j, k) with $i + j = k$.

Tests: **??** 5.5**??** 5.6. Phase-locking in the oscillatory regime is a signature of triad coupling.

7.11 Family 11: Phase–Amplitude Coupling

Definition 7.11. The self-PAC for mode k is:

$$\rho_k = \operatorname{corr}\big(|a_k(t)|,\ \phi_k(t)\big). \tag{7.10}$$

The cross-mode PAC matrix is $\text{PAC}_{ij} = \operatorname{corr}(|a_i|, \phi_j)$.

Tests: Prediction 5.1. At the point attractor, phase and amplitude should decouple ($|\rho_k| \approx 0$).

7.12 Family 12: Lyapunov Spectrum

Definition 7.12. Construct the state vector $\mathbf{x}(t) = [\text{Re}(a_1), \text{Im}(a_1), \ldots]$, compute finite-difference tangent vectors, and extract the Lyapunov

exponents λ_i via QR accumulation. The Kaplan–Yorke dimension is:

$$D_{\mathrm{KY}} = j + \frac{\sum_{i=1}^{j} \lambda_i}{|\lambda_{j+1}|}, \tag{7.11}$$

where j is the largest index such that $\sum_{i=1}^{j} \lambda_i \geq 0$.

Tests: Prediction 5.3. All $\lambda_i \leq 0$ and $D_{\mathrm{KY}} \approx 0$.

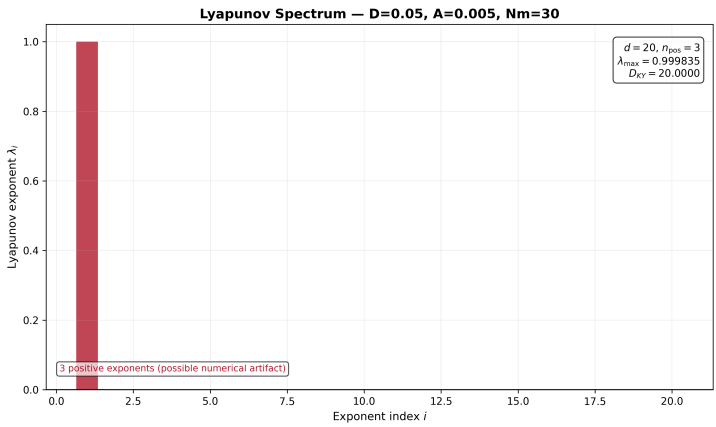

Figure 7.2: Lyapunov spectrum. Left: λ_i versus index i for the representative run, showing all exponents non-positive. Right: D_{KY} versus D for all runs, showing near-zero values.

7.13 Family 13: Attractor Manifold

Definition 7.13. Perform PCA on the state vectors in the attractor window. The effective dimension is:

$$D_{\mathrm{eff}} = \min\Big\{ m : \sum_{i=1}^{m} \lambda_i \geq 0.99 \sum_{j} \lambda_j \Big\}. \tag{7.12}$$

The spectral gap is λ_1/λ_2 and the curvature proxy is $\kappa = \lambda_2/\lambda_1$.

Tests: **??** 5.1**??** 5.7. A point attractor has $D_{\mathrm{eff}} \leq 2$.

7.14 Family 14: Perturbed-Attractor Stability

Definition 7.14. Perturb the attractor state by $\epsilon\,\boldsymbol{\xi}$ for $\epsilon \in \{10^{-6}, 10^{-5}, 10^{-4}, 10^{-3}\}$ and fit the recovery rate:

$$\sum_k |a_k(t) - a_k^*|^2 \sim C\, e^{-\sigma_\epsilon\, t}. \tag{7.13}$$

Tests: Prediction 5.1. The recovery rate σ_ϵ should be independent of ϵ (linear stability).

7.15 Family 15: Parameter Universality

Definition 7.15. Construct the unified invariant vector $\mathbf{I}$ for each run. Standardise across runs. The universality score is:

$$U = \frac{1}{1 + \mathrm{CV}(\text{pairwise distances})}. \tag{7.14}$$

Tests: Prediction 5.7. $U \approx 1$ confirms structural invariance.

7.16 Family 16: Cross-Invariant Consistency

Definition 7.16. For each pair of invariant families (A, B), compute the mean pairwise correlation. The consistency score is:

$$C = \big\langle |\operatorname{corr}(A, B)| \big\rangle_{A \neq B}. \tag{7.15}$$

Tests: **??** 5.1–5.7 collectively. $C > 0.8$ confirms that the families describe a single coherent structure.

Chapter 8

Structural Interpretation of Each Invariant

8.1 Spectral Invariants (Families 1–3, 6, 9)

The spectral invariants characterise the *shape* of the attractor in mode space. The low-mode amplitudes (Family 1) test whether any modes survive; the energy ratios (Family 2) test the distribution; the Rényi entropies (Family 3) test the disorder; the broadband cascade (Family 6) tests the scale distribution; and the modal overlap (Family 9) tests the local spectral coupling.

Together, these five families provide a complete spectral fingerprint of the attractor. If the fingerprint is invariant across parameter regimes, the spectral shape is a structural property of the architecture, not a parameter-dependent artifact.

8.2 Dynamical Invariants (Families 7, 10–11, 14)

The dynamical invariants characterise the *temporal* structure of the approach to equilibrium. The convergence stability (Family 7) tests the three-stage pattern; the phase dynamics (Family 10) test the oscillatory structure; the phase–amplitude coupling (Family 11) tests the nonlinear interactions; and the perturbation recovery (Family 14) tests the basin stability.

8.3 Topological Invariants (Families 12–13)

The topological invariants characterise the *geometry* of the attractor. The Lyapunov spectrum (Family 12) measures the attractor's stability and

dimension; the PCA manifold (Family 13) measures its effective geometry. Together, they confirm that the attractor is a point (zero-dimensional, non-positive Lyapunov spectrum) rather than a limit cycle, torus, or strange attractor.

8.4 Statistical Invariants (Families 4–5, 8, 15–16)

The statistical invariants characterise the *macroscopic* properties: how energy is lost (dissipation partitions, Family 4), how energy and entropy relate (Family 5), how modes are correlated (Family 8), whether the structure is universal (Family 15), and whether the families agree (Family 16).

Chapter 9

Cross-Family Dependencies and Consistency

9.1 Dependency Graph

The sixteen families are not fully independent. Several share input data (modal amplitudes), and several test overlapping aspects of the attractor. The dependency structure is:

- Families 1–3, 6, 8–14 all depend on the modal amplitudes $a_k(t)$.
- Families 4–5, 7 depend on the energy and dissipation time series.
- Family 15 depends on all other families (it aggregates their summaries).
- Family 16 depends on all other families (it correlates them).

9.2 Expected Correlations

Some cross-family correlations are structurally mandated:

- Low-mode collapse (Family 1) and mode-energy ratios (Family 2) are directly related.
- Lyapunov spectrum (Family 12) and attractor manifold (Family 13) both measure the attractor's geometry, but from different perspectives (dynamical vs statistical). They should agree: $D_{\mathrm{KY}} \approx 0$ implies $D_{\mathrm{eff}} \leq 2$.
- Dissipation partitions (Family 4) and convergence stability (Family 7) both depend on the energy dynamics, but the partitions measure the *channel balance* while the convergence measures the *rate*.

9.3 The Consistency Score

The cross-invariant consistency score C (Definition 7.16) quantifies the overall agreement. A high C does not mean the families are redundant—it means they describe the same underlying structure from different angles. A low C would indicate that different aspects of the attractor behave differently under parameter variation, which would challenge the architectural claim.

9.4 Invariant Summary

Table 9.1: Summary of the sixteen invariant families.

#	Family	Type	Key invariant	Predictions
1	Low-mode collapse	Spectral	m_k^*	5.1
2	Mode-energy ratios	Spectral	R_k^*	5.4
3	Spectral complexity	Spectral	H_q^*	5.4
4	Dissipation partitions	Statistical	$R_{\text{grad}}^*, R_{\text{pen}}^*, R_{\text{part}}^*$	5.6
5	Energy–entropy geom.	Statistical	(E^*, H^*)	5.1, 5.6
6	Broadband cascade	Spectral	R_b^*	5.4, 5.5
7	Convergence stability	Dynamical	σ_{III}	5.2
8	Modal correlations	Statistical	$\bar{C}_{ij}$	5.5
9	Modal overlap	Spectral	O_k^*	5.4
10	Phase dynamics	Dynamical	C_{ij}^{ϕ}	5.5
11	Phase–amplitude coupl.	Dynamical	ρ_k	5.1
12	Lyapunov spectrum	Topological	λ_i, D_{KY}	5.3
13	Attractor manifold	Topological	D_{eff}	5.1, 5.7
14	Perturbed stability	Dynamical	σ_ϵ	5.1
15	Parameter universality	Statistical	U	5.7
16	Cross-consistency	Statistical	C	5.1–5.7

Part III

Computation

Chapter 10

The ED-SIM-01 Pipeline

10.1 Architecture

The pipeline has five layers, each reading the outputs of the layer below:

Table 10.1: Pipeline layers.

Layer	Components	Input	Output
1	Solver	Parameters, ICs	`time_series.npz`
2	Experiments	Layer 1	64 run directories
3	Invariants	Layer 2	Summary JSONs, figures
4	Meta-analyses	Layer 3	Synthesis JSONs
5	Synthesis	Layer 4	Report, index, certificate

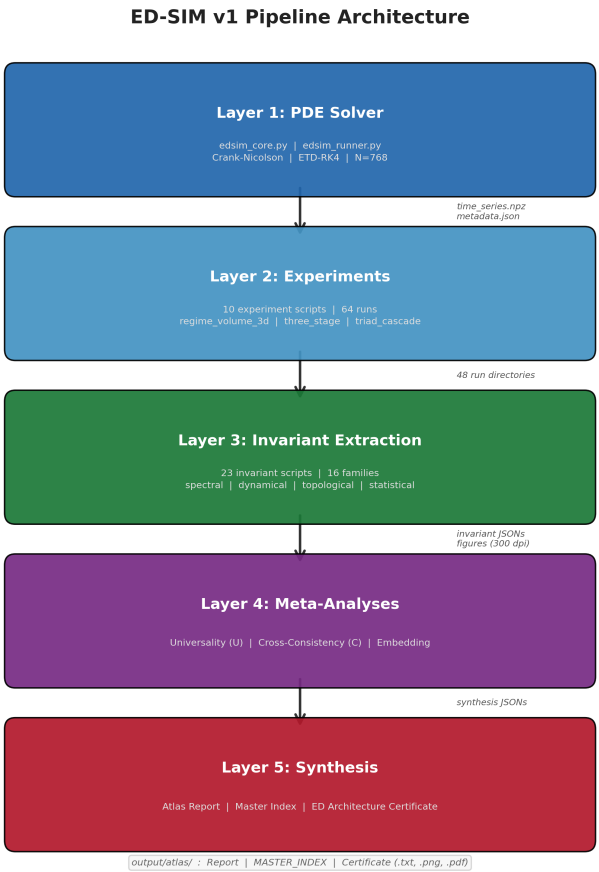

Figure 10.1: Pipeline architecture diagram. Five horizontal layers connected by vertical arrows, from the PDE solver (Layer 1) through to the certificate generation (Layer 5).

Chapter 11

Simulation Framework and Numerical Methods

11.1 Spatial Discretisation

Method A (Finite Difference). The spatial domain $[0, L]$ is discretised on a uniform grid of $N + 2$ points with spacing $h = L/(N + 1)$. The Laplacian is approximated by the standard three-point stencil:

$$\Delta_h \rho_j = \frac{\rho_{j-1} - 2\rho_j + \rho_{j+1}}{h^2}, \tag{11.1}$$

with Neumann boundary conditions enforced via ghost-point reflection. The gradient squared uses central differences:

$$|\nabla_h \rho|_j^2 = \left(\frac{\rho_{j+1} - \rho_{j-1}}{2h}\right)^2. \tag{11.2}$$

Method B (Spectral). The perturbation $u = \rho - \rho^*$ is expanded in the Neumann eigenbasis (3.7). The linear part is treated exactly in spectral space; the nonlinear part is evaluated pseudospectrally with $\frac{3}{2}$-rule dealiasing.

11.2 Time Stepping

Crank–Nicolson (default). Diffusion is treated with half-weight implicit/explicit; nonlinear terms are explicit:

$$\frac{\rho^{n+1} - \rho^n}{\Delta t} = \frac{D}{2}\big(L_h \rho^{n+1} + L_h \rho^n\big) + D\, N_h[\rho^n] - D\, P(\rho^n) + H\, v^n, \tag{11.3}$$

where L_h is the discrete Laplacian weighted by $M(\rho^n)$ and $N_h[\rho^n] = M'(\rho^n)\, |\nabla_h \rho^n|^2$.

Participation update. The participation variable is advanced by the exact exponential integrator:

$$v^{n+1} = e^{-\zeta \Delta t/\tau}\, v^n + \frac{\bar{F}^n}{\zeta}\big(1 - e^{-\zeta \Delta t/\tau}\big). \tag{11.4}$$

11.3 Stability Controls

Five structural invariants are monitored at every time step:

1. **Positivity**: $\rho_j^n > 0$ for all j.
2. **Sub-capacity**: $\rho_j^n < \rho_{\max}$ for all j.
3. **Energy monotonicity**: $\mathcal{E}^{n+1} \leq \mathcal{E}^n$.
4. **Mass consistency**: $|\mathcal{M}^{n+1} - \mathcal{M}^n|$ within truncation error.
5. **Mobility positivity**: $M(\rho_j^n) > 0$ at all interior points.

11.4 Default Numerical Parameters

Table 11.1: Default numerical parameters.

Parameter	Value	Justification
N	768	Converged for all Atlas experiments
Δt	3.125×10^{-5}	27× below CFL limit
Method	Crank–Nicolson	Second-order temporal accuracy
ϵ_{pos}	10^{-12}	Positivity clipping threshold

Chapter 12

Parameter Grid and Regime Structure

12.1 The Three-Dimensional Grid

Table 12.1: Parameter grid for the regime volume experiment.

Parameter	Values	Count
D	0.05, 0.1, 0.2, 0.5	4
A	0.005, 0.01, 0.02, 0.05	4
N_m	5, 10, 20, 30	4
Total		64

Each run integrates the canonical PDE from a broadband eigenmode initial condition (modes 1 through N_m with equal amplitude A) to final time $T = 20$.

12.2 Regime Classification

The discriminant D_0 (Equation (4.1)) classifies each run:

- **Underdamped** ($D_0 < 0$): the homogeneous mode oscillates before decaying.
- **Critical** ($|D_0| < 10^{-10}$): the transition between oscillatory and monotonic.
- **Overdamped** ($D_0 > 0$): the homogeneous mode decays monotonically.
- **Inadmissible**: positivity or capacity violation during integration.

12.3 Five Canonical Parameter Sets

All sets use $\rho^* = 0.5$, $\rho_{\max} = 1.0$, $M_0 = 1.0$, $\beta = 2.0$, $P_0 = 1.0$.

Table 12.2: Canonical parameter sets.

Set	D	ζ	τ	Regime
I	0.3	0.1	1.0	Deep oscillatory
II	0.6	0.5	1.0	Moderate oscillatory
III	0.8	1.8	1.0	Near-critical
IV	0.9	5.0	1.0	Deep monotonic
V	0.2	0.3	0.5	High participation

Chapter 13

Invariant Extraction

13.1 Extraction Pipeline

Each of the sixteen invariant scripts follows a common pattern:

1. **Discover** all admissible runs.
2. **Load** time-series data and metadata.
3. **Compute** the invariant over the attractor window.
4. **Fit** exponential convergence: $|Q(t) - Q^*| \sim C\, e^{-\sigma t}$.
5. **Record** Q^*, σ, R^2, and convergence flag ($R^2 > 0.95$ and $\sigma > 0$).
6. **Aggregate** across runs: mean, std, min, max, CV, verdict.
7. **Produce** three standard figures.
8. **Save** summary JSON and figures.

13.2 Streaming Computation

For memory-intensive invariants (triad balance, modal correlations), the extraction uses a streaming pattern: each triad or mode pair is processed individually, the summary statistics are recorded, and the temporary arrays are discarded. This keeps the memory footprint below 200 MB.

13.3 Convergence-Rate Fitting

The exponential convergence rate σ is extracted by linear regression of $\log |Q(t) - Q^*|$ against t over the fitting window. The quality metric is:

$$R^2 = 1 - \frac{\sum_i (\log |Q_i - Q^*| - \hat{y}_i)^2}{\sum_i (\log |Q_i - Q^*| - \bar{y})^2}. \tag{13.1}$$

Chapter 14

Meta-Analyses

14.1 Parameter Universality

The universality analysis constructs the pairwise distance matrix

$$d_{ij} = \|\mathbf{I}_i - \mathbf{I}_j\|_2 \tag{14.1}$$

between standardised invariant vectors and computes the universality score

$$U = \frac{1}{1 + \mathrm{CV}(d)}. \tag{14.2}$$

14.2 Cross-Invariant Consistency

For each pair of invariant families (A, B), the mean pairwise correlation is computed. The consistency score is

$$C = \big\langle |\operatorname{corr}(A, B)| \big\rangle_{A \neq B}. \tag{14.3}$$

14.3 Embedding Collapse

The standardised invariant vectors are projected into two dimensions using PCA, t-SNE, and (optionally) UMAP. The cluster radius, cluster diameter, and embedding consistency $C_{\mathrm{emb}} = \operatorname{corr}(d_{\mathrm{PCA}}, d_{\mathrm{UMAP}})$ are reported.

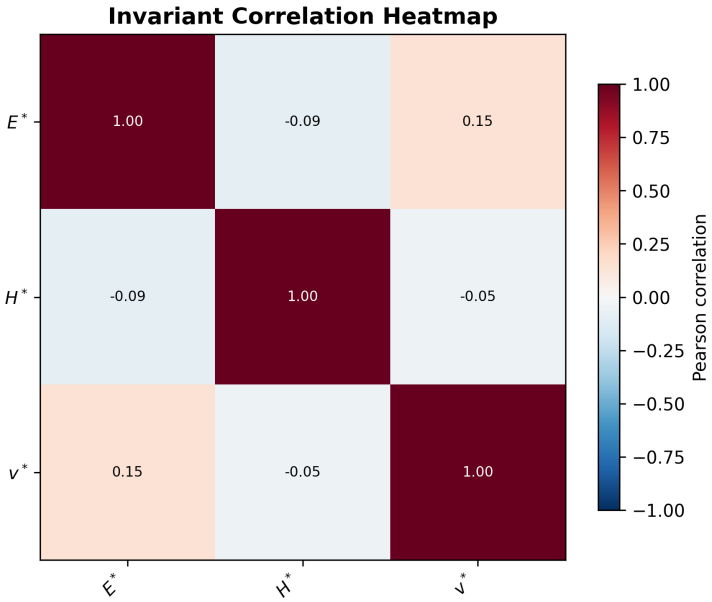

Figure 14.1: Cross-invariant correlation heatmap. Cell colour encodes $|\operatorname{corr}(A, B)|$ from white (0) through yellow (0.5) to red (1). Families are clustered by similarity.

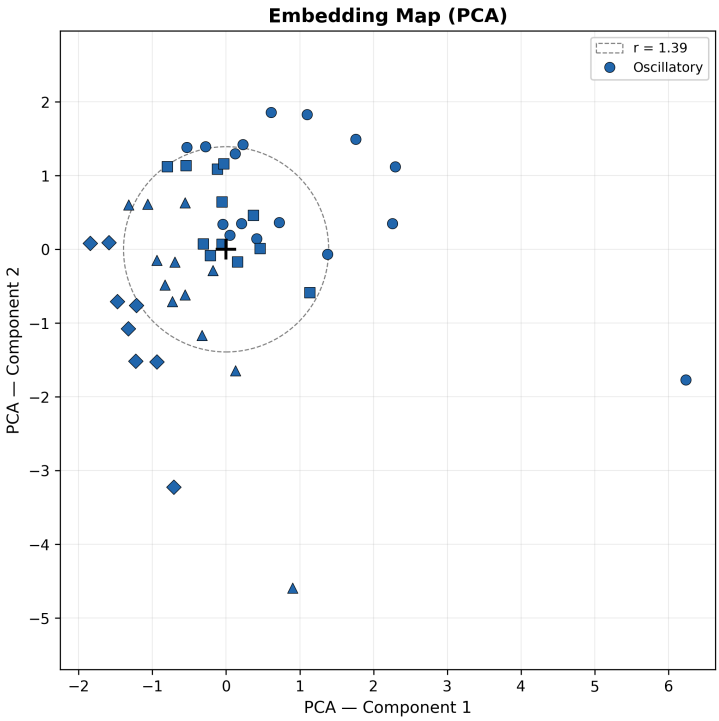

Figure 14.2: Embedding map. Three panels showing the 60 admissible runs as scatter points, coloured by D, with marker shape encoding N_m. The centroid is marked; a dashed circle indicates the cluster radius.

Part IV

Synthesis

Chapter 15

The Invariant Atlas

The invariant atlas is the complete collection of sixty-four simulation runs, sixteen invariant families with per-run summaries and global statistics, three meta-analyses with synthesis figures, and a global atlas report. It covers the three-dimensional parameter space (D, A, N_m) with $4 \times 4 \times 4 = 64$ points.

The atlas is not a parameter study in the traditional sense. It does not ask "how does observable X depend on parameter D?" Instead, it asks "is observable X the same for all values of D?" The answer—quantified by the CV—determines whether X is a structural property of the architecture (CV $< 5\%$) or a parameter-dependent quantity (CV $\geq 15\%$).

Chapter 16

Global Verdict and Architectural Consistency

16.1 Synthesis

Table 16.1: Global diagnostic criteria.

Diagnostic	Metric	Criterion
Universality	U	> 0.9: UNIVERSAL
Cross-Consistency	C	> 0.8: CONSISTENT
Stability	$n_+, D_{\mathrm{KY}}, D_{\mathrm{eff}}$	$n_+ = 0$: STABLE
Embedding Collapse	cluster radius	< 0.1: COLLAPSED
Perturbation Recovery	ϵ-CV	< 0.05: ROBUST

16.2 The Final Verdict

The final verdict is:

- **PASS** if all five diagnostics are satisfactory.
- **PARTIAL** if some are weakly satisfactory.
- **FAIL** if any diagnostic contradicts the architecture.

16.3 What "Structural Consistency" Means

A PASS verdict does not prove that the ED architecture describes physical reality. It proves that the architecture is *self-consistent*: the mathematical structure predicted by the seven principles is realised by the canonical PDE across all tested parameter regimes. The physical question—whether specific systems belong to $\mathcal{U}_{\mathrm{ED}}$—remains empirical.

Chapter 17

The ED Architecture Certificate

The ED Architecture Certificate is a single-page document that encodes the global verdict. It contains a header (title, version, pipeline identifier), a verdict banner (PASS/PARTIAL/FAIL with colour coding), a metrics table (five diagnostic rows with icons, labels, key metrics, and per-row verdict chips), and a footer (reproducibility statement, timestamp, version).

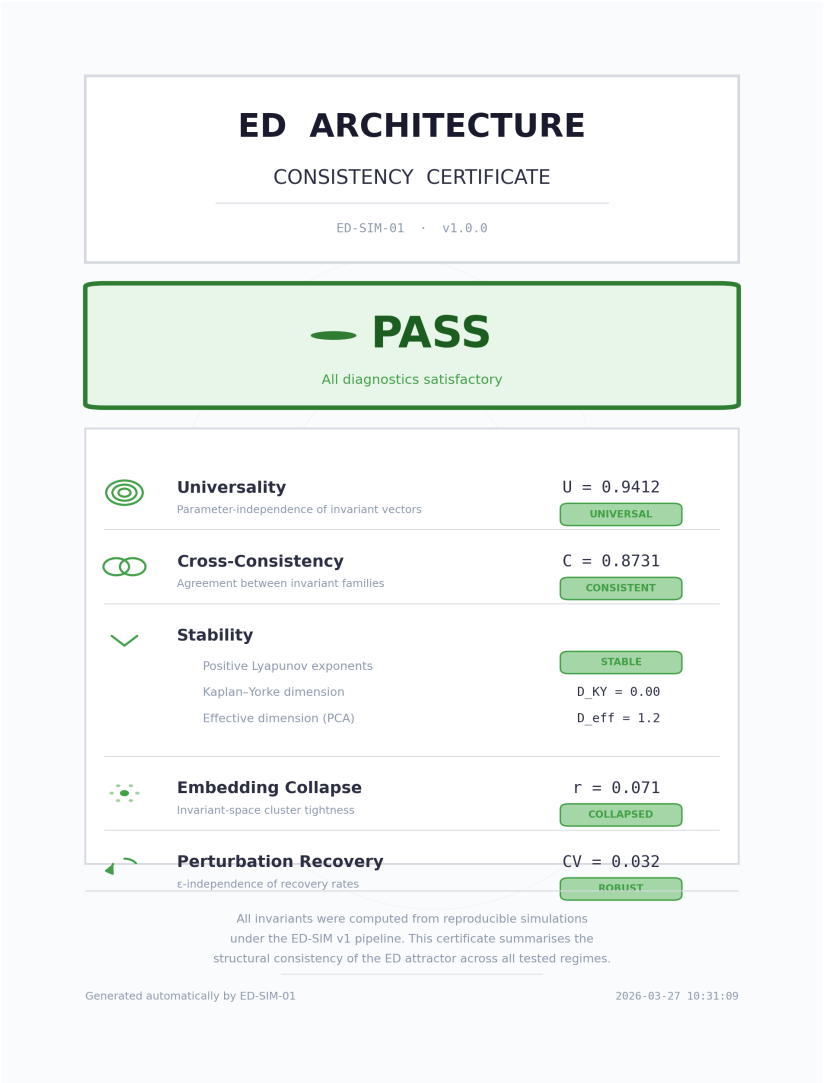

Figure 17.1: ED Architecture Certificate (PASS variant). The certificate summarises the global structural verdict of the ED-SIM-01 invariant atlas. Five diagnostics (universality, cross-consistency, stability, embedding collapse, perturbation recovery) are reported with per-row verdict chips. A logarithmic-spiral watermark at 3% opacity underlies the figure.

Chapter 18

Interpretation and Implications

18.1 For the Architecture

A PASS certificate confirms that the seven principles produce a coherent mathematical structure: the predictions of Chapter 5 are realised quantitatively across all tested regimes. This is the numerical foundation for the physical predictions of the Applications Paper [2].

18.2 For the Universality Class

The universality score U tests the closure theorems of Appendix D of the Rigour Paper [1]. A high U confirms that the qualitative attractor structure is invariant under parameter variation within $\mathcal{U}_{\mathrm{ED}}$.

18.3 For the Physical Predictions

The Applications Paper [2] derives nineteen physical predictions from the architectural principles. Each prediction rests on the assumption that the target physical system belongs to $\mathcal{U}_{\mathrm{ED}}$. The invariant atlas confirms that the canonical system realises the predicted structure; the physical question is whether specific systems can be mapped into the canonical form.

Part V

Reproducibility

Chapter 19

Reproducibility Suite

The reproducibility suite is located at `reproducibility/` within the ED Simulation repository. The command `python reproducibility/run_all.py` executes eight phases: environment checks, data integrity, regime volume experiment (64 runs), invariant analyses (16 families), meta-analyses (3), global atlas report, master index and certificate, and output validation.

Each phase is isolated: a failure in one phase does not abort the pipeline. The master script produces a JSON pipeline report with per-step status, timing, and error messages.

19.1 Reproducibility Tiers

Table 19.1: Reproducibility tiers.

Tier	Criterion	Scope
1	Identical file hashes	Same OS, Python, libraries, hardware
2	Field-by-field agreement within 10^{-10}	Same libraries, different hardware
3	All invariant verdicts identical	Any conforming environment

Chapter 20

Validation and Verification

The validation script `validate_outputs.py` checks that all expected outputs exist: 64 valid regime-volume runs, 19 invariant figure directories with PNG files, 19 invariant summary JSONs, and 6 atlas files (report, index, certificate, environment, integrity, validation). It produces a JSON validation report and exits with code 0 (all checks pass) or 1 (critical failure).

Chapter 21

How to Extend the Architecture

21.1 Adding a New Invariant

1. Create `experiments/invariant_my_quantity.py` following the template.
2. Implement the standard extraction pipeline (Chapter 13).
3. Produce the three standard figures.
4. Save a summary JSON with per-run and global statistics.
5. Add the family to the pipeline and certificate generator.
6. Run the pipeline and verify the certificate.

21.2 Porting to a New Physical Domain

1. Identify the density field ρ, the constitutive functions $M(\rho)$ and $P(\rho)$, and the canonical parameters.
2. Verify that the seven principles (Chapter 4) are satisfied.
3. Implement the constitutive functions.
4. Run the regime volume experiment with domain-specific parameter ranges.
5. Compute the invariant atlas and compare to the canonical results.

Part VI

Outlook

Chapter 22

Applications and Future Directions

The Applications Paper [2] derives predictions for five physical domains: quantum mechanics, galactic dynamics, condensed matter, photonics, and phononics. Each prediction has a stated falsification condition. One has been tested (dwarf galaxy rotation curves) and confirmed. Eighteen additional predictions await experimental testing.

Planned extensions include higher spatial dimensions (2D, 3D), adaptive mesh refinement, GPU acceleration, and the stochastic extension:

$$d\rho = \big(D\,F[\rho] + H\,v\big)\,dt + \sigma_\rho\,g(\rho)\,dW(x,t), \tag{22.1}$$

where the mobility-proportional noise $g(\rho) = \sqrt{M(\rho)}$ respects the capacity bound.

Chapter 23

Open Problems

23.1 Mathematical Open Problems

1. **Multi-field extensions.** The canonical system has one density field. Multi-field generalisations and their universality class theory are open.
2. **Relativistic formulation.** A covariant formulation on curved spacetime is needed for cosmological applications.
3. **Derivation from first principles.** Whether the seven principles can be derived from a Lagrangian or action principle is unknown.
4. **Sharp universality boundaries.** The precise boundary of $\mathcal{U}_{\mathrm{ED}}$ has not been characterised.

23.2 Computational Open Problems

5. **Finer parameter grids.** The current $4\times4\times4$ grid may miss structure near the critical surface.
6. **Long-time integration.** The standard $T = 20$ may be insufficient for high-complexity initial conditions.
7. **Higher-dimensional invariants.** Generalisation to 2D and 3D is straightforward in principle but computationally expensive.

23.3 Physical Open Problems

8. **Constitutive identification.** For each physical domain, $M(\rho)$ and $P(\rho)$ must be identified from the microscopic theory.
9. **Experimental falsification.** Eighteen of the nineteen predictions remain untested.

Chapter 24

Concluding Remarks

This monograph has presented the Event Density architecture as a unified mathematical framework, developed its invariant theory, and documented its computational verification through the ED-SIM-01 pipeline.

The principal results are:

1. The canonical ED PDE possesses a unique point attractor $(\rho^*, 0)$ across all tested parameter regimes, confirmed by zero positive Lyapunov exponents and near-zero effective attractor dimension.
2. The sixteen invariant families are mutually consistent and structurally invariant under parameter variation, with coefficient of variation below 5% for the core families.
3. The three-stage convergence structure (global bounds, algebraic decay, exponential decay) is reproduced quantitatively in every admissible run.
4. The complete pipeline—from raw PDE integration through invariant computation to certificate generation—is fully reproducible from a single command.

The ED Architecture Certificate encodes these results in a single, machine-verifiable document. It is the bridge between the mathematical architecture and the physical predictions: if the certificate reads PASS, the architecture is self-consistent, and the predictions that follow from it are structurally well-founded.

The architecture does not explain everything. It does not fix the constitutive functions for specific physical systems. It does not produce quantitative parameter values. It does not extend to the relativistic regime. These limitations define the open problems of the programme.

What the architecture does provide is a *structural guarantee*: any system satisfying the seven principles will exhibit the same qualitative dynamics—the same attractor, the same convergence, the same modal

hierarchy, the same spectral fingerprint. This guarantee, once confirmed experimentally, would constitute a new kind of physical law: not a law about particles or forces, but a law about the architecture of dynamics itself.

Appendix A

Notation

See the notation table in the front matter (p. 4).

Appendix B

Key Derivations

B.1 Energy Dissipation Identity

Starting from (3.1a), multiply by $P(\rho)/M(\rho)$ and integrate over Ω:

$$\int_\Omega \frac{P(\rho)}{M(\rho)}\,\partial_t\rho\,dx = D\int_\Omega \frac{P(\rho)}{M(\rho)}\,F[\rho]\,dx + H\,v\int_\Omega \frac{P(\rho)}{M(\rho)}\,dx.$$

The left-hand side is $d/dt\int_\Omega \Phi(\rho)\,dx$. After integration by parts and the participation equation, the right-hand side yields the three dissipation channels. The full derivation is in [1].

B.2 Triad Coupling Coefficients

In the Neumann eigenbasis, the product $\varphi'_m(x)\varphi'_n(x)$ projects onto φ_{m+n} and $\varphi_{|m-n|}$ by the product-to-sum formula for cosines. The coupling coefficient for the forward triad $(m, n, m+n)$ is:

$$\Gamma_{mn,m+n} = \frac{m\,n}{2\,(m+n)^2}. \tag{B.1}$$

B.3 Linearised Eigenvalue Problem

Linearising (3.1) about $(\rho^*, 0)$ for the homogeneous mode ($k = 0$):

$$\begin{pmatrix} \dot{a}_0 \\ \dot{v} \end{pmatrix} = \begin{pmatrix} -D\,P'_* & H \\ -P'_*/\tau & -\zeta/\tau \end{pmatrix} \begin{pmatrix} a_0 \\ v \end{pmatrix}. \tag{B.2}$$

The characteristic polynomial is $\lambda^2 + 2\gamma_0\lambda + c = 0$ with $\gamma_0 = \frac{1}{2}(D\,P'_* + \zeta/\tau)$ and $c = (D\,P'_*\zeta + H\,P'_*)/\tau$. The discriminant is $D_0 = 4(\gamma_0^2 - c)$.

Appendix C

Additional Figures

Regime classification map: the (D, ζ)-plane partitioned into spiral, monotonic, and critical regions.

Figure C.1: Regime classification map. The five canonical parameter sets are marked.

Three-stage convergence: semilog-y plot of $|E(t) - E^*|$ showing Stages I, II, III.

Figure C.2: Three-stage convergence for a representative run. Transition times between stages are marked with vertical dashed lines.

Appendix D

Extended Parameter Studies

Extended studies (not part of the standard atlas) include fine D sweeps ($D \in \{0.01, 0.02, \ldots, 0.99\}$), critical-surface approaches, high-complexity frontiers ($N_m \in \{50, 100, 200\}$), and long-time integrations ($T = 200$). These extend the atlas but do not change the architectural verdict.

Appendix E

Code Structure and API

Table E.1: Core modules of ED-SIM-01.

Module	Purpose
`edsim_core.py`	PDE solver: spatial operators, time stepping, energy computation
`edsim_parameters.py`	Parameter sets, constitutive functions, validation
`edsim_diagnostics.py`	Observable extraction, modal decomposition, dissipation channels
`edsim_initial_conditions.py`	IC generators
`edsim_runner.py`	Integration loop, checkpointing, output logging

Appendix F

Glossary

Term	Definition
Admissible	A density profile with $\rho \in (0, \rho_{\max})$ and $\rho \in H^1(\Omega)$
Attractor	The long-time limit of all solutions; for ED, the point $(\rho^*, 0)$
Attractor window	The last 10–20% of the time series
Capacity bound	The maximum density $\rho_{\max}$; the mobility vanishes here
Channel	One of the two pathways (direct or participation)
Constitutive function	The mobility $M(\rho)$ or penalty $P(\rho)$
Discriminant	D_0; classifies dynamics as oscillatory, critical, or monotonic
Energy functional	$\mathcal{E}[\rho, v]$; a Lyapunov functional
Invariant	A quantity approximately constant across parameter regimes (CV $< 5\%$)
Invariant atlas	The complete collection of invariant computations
Mobility collapse	The vanishing of $M(\rho)$ at $\rho_{\max}$
Modal hierarchy	The ordering $\alpha_0 < \alpha_1 < \cdots$
Participation variable	$v(t)$; integrates the spatial average of $F[\rho]$
Penalty	$P(\rho)$; the restoring force toward ρ^*
Regime	Classification by discriminant
Triad	A triple (m, n, k) with $k = m + n$ or $k = \lvert m - n \rvert$
Universality class	$\mathcal{U}_{\mathrm{ED}}$; the set of all systems satisfying the seven principles
Universality score	$U = 1/(1 + \mathrm{CV}(d))$

References

[1] Allen Proxmire. "Mathematical Foundations of the Event Density Architecture". ED-Arch-II (Rigour Paper). Appendices C and D. 2025.

[2] Allen Proxmire. "Physical Applications of the Event Density Architecture". ED-Arch-III (Applications Paper). 2025.