The Architectural Canon of the Unified ED Cosmology PDE

Seven Irreducible Principles Defining the Event-Density Architecture

Allen Proxmire
March 2026

Abstract

The Event-Density (ED) cosmology PDE exhibits a rich, multi-layered architectural structure encompassing motifs, laws, geometry, field theory, energetics, causality, information, symmetry, and invariants. While the ED-Arch series (ED-Arch-01 through ED-Arch-11) decomposed and reconstructed this structure, it did not identify its minimal foundation. This paper extracts the Architectural Canon: the smallest set of structural principles from which the entire ED architecture follows. We show that exactly seven principles — operator structure, channel complementarity, penalty equilibrium, mobility capacity bound, participation feedback loop, damping discriminant, and nonlinear triad coupling — are jointly sufficient and individually necessary to generate all architectural layers. These principles define the identity conditions for ED: any PDE satisfying them is architecturally ED, and any PDE violating even one is not. The Canon also defines the canonical equivalence class of ED systems and yields a canonical normal form for the architecture. By establishing the irreducible basis of the ED ontology, the Architectural Canon completes the foundational arc of the ED-Arch series and provides the structural anchor for future work in architectural applications and cosmological interpretation.

1. Introduction

1.1 Context and Motivation

The unified ED cosmology PDE is a structurally rich dynamical system whose behavior cannot be understood from its equation alone. Over the course of the ED-Arch series (ED-Arch-01 through ED-Arch-11), the architecture underlying this PDE was progressively uncovered: first through motifs, then through laws, geometry, field theory, energetics, causality, information, symmetry, and invariants. Each layer revealed a deeper structural regularity, and each layer constrained the next.

By the end of ED-Arch-11, the architecture stood as a complete, multi-layered ontology. Yet this ontology, while coherent, was not minimal. Many of its features were consequences of deeper structural choices. The purpose of the present paper is to identify those choices — the irreducible principles from which the entire architecture follows.

This document extracts the Architectural Canon: the smallest set of structural principles that uniquely define the ED architecture and generate all of its motifs, laws, geometric objects, field components, energy functionals, causal channels, information transformations, symmetries, and invariants. These principles are not derived; they are the axioms from which everything else descends.

1.2 What an Architectural Canon Is

In the ED context, an architectural canon is defined by three criteria:

(i) Minimal Set

The canon is the smallest possible collection of independent structural principles from which the entire ED architecture can be derived. If any principle is removed, some part of the architecture becomes underivable. If any principle is added, it is redundant.

(ii) Irreducible Core
None of the canonical principles can be derived from the others. Each stands as a structural axiom. Together they form a basis — in the algebraic sense — for the architecture's full content. The canon contains no internal redundancy.

(iii) Identity Conditions
The canon defines what it means for a PDE to be architecturally ED.
Any PDE satisfying all canonical principles is ED, regardless of its specific functional forms.
Any PDE violating even one principle is not ED, regardless of superficial similarity.
The canon is the architecture's fingerprint.

1.3 What the Canon Is Not
The canon must be distinguished from three other structural concepts developed earlier in the ED-Arch series:

Laws (ED-Arch-03)
The nine architectural laws describe universal constraints within the ED architecture.
They are derived statements.
The canon explains why the laws hold.

Invariants (ED-Arch-10)
The nine irreducible invariants and seven invariant relations describe what remains unchanged under architectural transformations.
They are consequences.
The canon explains why these invariants exist.

Motifs (ED-Arch-02)
The eight motifs are recurring structural patterns in the dynamics.
They are manifestations.
The canon explains why these motifs appear and why no others do.

The hierarchy is:

```
CANON → LAWS → MOTIFS → BEHAVIORS → FEATURES
                                        ↘ INVARIANTS
                                        ↘ GEOMETRY
                                        ↘ ALL OTHER LAYERS
```

The canon sits at the apex. Everything else descends from it.

1.4 Why the Canon Is Needed
The nine-layer architecture developed in ED-Arch-02 through ED-Arch-11 is rich but contains structural redundancy. Many features at higher layers are consequences of deeper choices:

- The laws imply the geometry.
- The energetics imply the causal structure.
- The symmetries imply the invariants.

- The motifs arise from combinations of laws and geometry.

Working within the full nine layers, it is not immediately clear which structural elements are fundamental and which are derivative.

The Canon resolves this by identifying the irreducible foundation.
It answers the question:

> **What is the smallest set of structural commitments required to specify the ED architecture completely?**

The answer is seven.
Seven independent principles generate the entire ontology.

This matters for three reasons:

Identity
The Canon tells us what ED is.
Not eight motifs, nine laws, or nine invariants — but seven principles.

Universality
Any PDE satisfying the same seven principles will exhibit the same motifs, obey the same laws, and possess the same invariants, even if its functional forms differ.
The Canon explains why diverse systems can be architecturally identical.

Fragility
Violating any one principle breaks the architecture in specific, predictable ways.
The Canon identifies the architecture's minimal attack surface.

2. The Seven Canonical Principles

The ED architecture is generated by exactly seven irreducible structural principles. Each principle is independent, none can be derived from the others, and together they form the minimal basis from which the entire nine-layer ED ontology follows. This section presents each principle in its canonical form: its statement, its irreducibility, the architectural layers it generates, and the structural consequences of its removal.

2.1 Canonical Principle 1 — Operator Structure

Statement
The density field evolves under a nonlinear operator of the form

$$F[\rho] = M(\rho)\nabla^2\rho + M'(\rho)|\nabla\rho|^2 - P(\rho)$$

comprising curvature-dependent diffusion, gradient-squared nonlinearity, and a density-dependent restoring force.

Why It Is Irreducible

- No other principle specifies the form of the operator.

- Channel complementarity determines how the operator enters the dynamics, not what it is.
- Penalty equilibrium constrains the zero of P, not its existence.
- The mobility bound constrains M at ϱmax, not its role.

Thus the operator structure is a primitive architectural choice.

Architectural Layers Generated

- Motifs: Saddle, Dissipator, Nonlinear Triad
- Laws: Modal Hierarchy, Harmonic Protection, Transient Structure
- Geometry: Modal Decay Funnel, Triad Locus
- Fields: Operator field
- Energetics: VM, VN, D-channel dissipation
- Causality: Spatial and spectral propagation
- Information: Routing, mixing, compression
- Symmetry: Spatial translation symmetry
- Invariants: Gradient sign, modal ordering, triad ratio

What Breaks If Removed
No spatial dynamics, no diffusion, no harmonic generation, no triad, no saddle, no modal hierarchy.
The architecture collapses to uncoupled ODEs.

2.2 Canonical Principle 2 — Channel Complementarity

Statement
Density evolves through two complementary channels:

$$\dot{\rho} = D\,F[\rho] + H\,v, \qquad D + H = 1.$$

The direct channel applies the operator immediately; the mediated channel applies it through memory.

Why It Is Irreducible
Nothing else in the architecture requires two channels or a conservation constraint.
The principle is the architectural decision that density evolution is a convex combination of direct and mediated influence.

Architectural Layers Generated

- Motifs: Boundary, Memory Loop, Dissipator
- Laws: Channel Complementarity, Critical Damping Boundary
- Geometry: Boundary Surface
- Fields: Dual-channel coupling
- Energetics: Channel-specific dissipation
- Causality: Spatial vs temporal influence
- Information: Encoding regime switching
- Symmetry: Multiplicative channel symmetry
- Invariants: $D+H = 1$

What Breaks If Removed
No principled combination of direct and mediated influence; no boundary motif; no regime structure; no architectural organization of parameter space.

2.3 Canonical Principle 3 — Penalty Equilibrium

Statement
The penalty function $P(\varrho)$ has exactly one zero at $\varrho = \varrho^{*}$, with P strictly increasing.
This equilibrium is unique and globally attracting.

Why It Is Irreducible
No other principle enforces uniqueness or monotonicity of the equilibrium.
This is the architectural decision that the system has a single target state.

Architectural Layers Generated

- Motifs: Manifold
- Laws: Unique Attraction, One-Dimensional Relaxation, Proportional Stability
- Geometry: Attractor point, stable manifold
- Fields: Restoring-force component
- Energetics: Penalty potential *VP*
- Causality: Direction of becoming
- Information: Erasure pipeline
- Symmetry: Broken density-translation symmetry
- Invariants: Unique attractor, monotone energy decrease

What Breaks If Removed
Multiple equilibria, limit cycles, or unbounded growth.
No manifold, no direction of becoming, no information erasure.

2.4 Canonical Principle 4 — Mobility Capacity Bound

Statement
The mobility satisfies $M(\varrho_{max}) = 0$ with $M(\varrho) > 0$ for $\varrho < \varrho_{max}$.
Density is bounded above, and diffusive transport collapses near the bound.

Why It Is Irreducible
No other principle imposes an upper bound on density or requires mobility collapse.
This is the architectural decision that the system has finite capacity.

Architectural Layers Generated

- Motifs: Horizon
- Laws: Finite Capacity
- Geometry: Horizon Surface
- Fields: Mobility field
- Energetics: D-channel blockade

- Causality: Diffusive causal barrier
- Information: Truncation
- Symmetry: Broken amplitude scaling
- Invariants: Capacity bound, horizon threshold

What Breaks If Removed
Unbounded density, no horizon, no finite capacity law, no kinetic saturation.

2.5 Canonical Principle 5 — Participation Feedback Loop

Statement

The participation variable evolves as

$$\dot{v} = \frac{1}{\tau}(F[\rho] - \zeta v)$$

integrating the operator with exponential decay and feeding back into density evolution.

Why It Is Irreducible
No other principle introduces memory or a secondary variable.
This is the architectural decision that the system accumulates and re-injects operator output.

Architectural Layers Generated

- Motifs: Spiral, Memory Loop
- Laws: Critical Damping Boundary
- Geometry: Spiral Sheet
- Fields: Participation field
- Energetics: V_v, ζ-channel dissipation
- Causality: Temporal propagation
- Information: Storage, delay, feedback
- Symmetry: Temporal reparameterization
- Invariants: Clockwise winding, loop topology

What Breaks If Removed
No memory, no oscillation, no spiral, no temporal causality, no information storage.
The architecture collapses to a single-variable diffusion.

2.6 Canonical Principle 6 — Damping Discriminant

Statement

Flow type is determined by the discriminant

$$D + 2\zeta.$$

Under-damped when $D+2\zeta < 1$, over-damped when $D+2\zeta > 1$, with a sharp transition at 1.

Why It Is Irreducible
Neither channel complementarity nor the participation loop alone predicts the discriminant or its threshold. The discriminant is an emergent structural boundary that cannot be derived from any single principle.

Architectural Layers Generated

- Motifs: Boundary
- Laws: Critical Damping Boundary
- Geometry: Boundary Surface, Spiral Sheet vs Monotonic Cone
- Fields: Regime-dependent topology
- Energetics: Dissipation profile transition
- Causality: Regime switch
- Information: Encoding switch
- Symmetry: Linear boundary symmetry
- Invariants: Linear discriminant

What Breaks If Removed
No distinction between oscillatory and monotonic regimes; no boundary motif; no partition of parameter space.

2.7 Canonical Principle 7 — Nonlinear Triad Coupling

Statement

The nonlinear term $M'(\varrho)|\nabla\varrho|^2$ generates higher-harmonic content, especially $k = 3$ from $k = 1$, with an invariant amplitude ratio under multiplicative perturbations.

Why It Is Irreducible
The operator structure asserts the existence of the term; the triad principle asserts its spectral consequences. The locked harmonic ratio is not implied by any other principle.

Architectural Layers Generated

- Motifs: Nonlinear Triad
- Laws: Harmonic Protection
- Geometry: Triad Locus
- Fields: Nonlinear coupling
- Energetics: Nonlinear interaction energy
- Causality: Spectral propagation
- Information: Mixing
- Symmetry: Broken spectral permutation symmetry
- Invariants: Locked k_3/k_1 ratio

What Breaks If Removed
No harmonic generation, no spectral fingerprint, no triad motif, no harmonic protection law.

3. Canonical Dependency Graph

The seven canonical principles do not merely constrain the ED architecture — they generate it. Every motif, law, geometric object, field component, energetic channel, causal rule, information transformation, symmetry, and invariant arises as a direct or indirect consequence of these principles. This section traces those dependencies explicitly, demonstrating that the Canon is both complete (it generates all architectural layers) and minimal (no layer requires additional assumptions).

3.1 Principles → Motifs

Each of the eight architectural motifs is generated by one or more canonical principles. No motif requires assumptions beyond the Canon.

Motif	Generated By		
Saddle	P1 (operator structure: ∇^2 sign reversal)		
Spiral	P5 (participation loop) + P6 (damping discriminant < 1)		
Horizon	P4 (mobility capacity bound)		
Manifold	P3 (unique equilibrium) + P1 (modal spectral gap)		
Boundary	P6 ($D + 2\zeta = 1$ threshold)		
Memory Loop	P5 (feedback loop) + P2 (H·v coupling)		
Dissipator	P1 (diffusion) + P2 (direct channel D·F)		
Nonlinear Triad	P7 (M′	$\nabla\rho$	^{2} coupling)

Coverage: All eight motifs are fully accounted for.
Implication: The Canon generates the architecture's entire pattern vocabulary.

3.2 Principles → Laws

Each of the nine architectural laws is a direct consequence of one or more canonical principles.

Law	Generated By
Unique Attraction	P3
One-Dimensional Relaxation	P3 + P1
Proportional Stability	P3
Critical Damping Boundary	P6 + P2 + P5
Finite Capacity	P4
Harmonic Protection	P7
Modal Hierarchy	P1
Channel Complementarity	P2
Transient Structure	P1 + P3 + P5

Coverage: All nine laws are grounded in the Canon.
Implication: The Canon is the source of all architectural constraints.

3.3 Principles → Geometric Objects

The eight geometric objects of ED-Arch-04 arise from the Canon as follows:

Geometric Object	Generated By
Attractor Point	P3
Stable Manifold Curve	P3 + P1
Spiral Sheet	P5 + P6
Monotonic Cone	P1 + P6
Horizon Surface	P4
Boundary Surface	P6
Modal Decay Funnel	P1
Nonlinear Triad Locus	P7

Coverage: All geometric structures are derivable.
Implication: Geometry is not assumed — it emerges from the Canon.

3.4 Principles → Field Components

The three coupled fields and their local encodings arise from the Canon:

Density Field (ρ)
Generated by:

- P1 (operator acts on ρ)
- P3 (restoring force)
- P4 (capacity bound)

Participation Field (v)
Generated by:

- P5 (integration of F[ρ])
- P2 (H·v coupling into ρ)

Operator Field (F[ρ])
Generated by:

- P1 (definition of F)
- P7 (nonlinear coupling structure)

Implication: The Canon generates the full field theory.

3.5 Principles → Energetics

The five energy functionals and three dissipation channels arise from:

Energetic Component	Generated By
Penalty Energy V_P	P3
Mobility Energy V_M	P1
Nonlinear Energy V_N	P7
Memory Energy V_v	P5
Total Energy E	P1 + P3 + P5
D-channel Dissipation	P1 + P2
ζ-channel Dissipation	P5
P-channel Dissipation	P3

Implication: Energetics is not an add-on — it is structurally forced.

3.6 Principles → Causality

The three causal channels and five propagation rules arise from:

Causal Component	Generated By
Spatial Causality	P1
Temporal Causality	P5
Spectral Causality	P7
Propagation Rules	P1 + P2 + P5 + P7

Implication: Causality is a direct expression of the Canon.

3.7 Principles → Information

The nine-stage information pipeline (routing → mixing → compression → storage → feedback → delay → truncation → collapse → erasure) arises from:

- P1 (routing, mixing, compression)
- P5 (storage, feedback, delay)
- P4 (truncation)
- P3 (collapse, erasure)

Implication: Information processing is not emergent — it is architecturally mandated.

3.8 Principles → Symmetry and Symmetry Breaking

Preserved symmetries arise from:

- P1 (spatial translation, sign inversion)
- P2 (multiplicative channel symmetry)
- P5 (temporal reparameterization)

Broken symmetries arise from:

- P3 (density translation)
- P4 (amplitude scaling)
- P6 (flow-type symmetry)
- P7 (spectral permutation)

Implication: The Canon determines both symmetry and its breaking.

3.9 Principles → Invariants

The nine irreducible invariants and seven invariant relations arise from:

- P1 (modal ordering, gradient sign)
- P2 (D+H = 1)
- P3 (unique attractor, monotone energy decrease)
- P4 (capacity bound)
- P5 (winding direction, loop topology)
- P6 (linear discriminant)
- P7 (locked harmonic ratio)

Implication: Invariants are not empirical — they are architectural.

4. Canonical Minimality

The Architectural Canon is not merely a convenient summary of the ED architecture — it is the minimal and irreducible foundation from which the entire ontology follows. This section demonstrates that the seven canonical principles are (i) independent, (ii) jointly sufficient, and (iii) individually necessary. No principle can be removed without breaking the architecture, and no principle can be derived from the others.

4.1 Independence of the Seven Principles

Each canonical principle introduces a structural commitment that none of the others imply. Their independence can be demonstrated by examining the architectural role of each principle and showing that no combination of the remaining six can reproduce it.

P1 — Operator Structure

No other principle specifies the form of the operator, the existence of curvature-dependent diffusion, or the gradient-squared nonlinearity. Without P1, the PDE has no spatial dynamics, no harmonic generation, and no spectral structure.

P2 — Channel Complementarity

No other principle introduces two complementary channels or the conservation constraint $D+H = 1$. Without P2, there is no principled way to combine direct and mediated influence.

P3 — Penalty Equilibrium

No other principle enforces a unique, monotone equilibrium. Without P3, the system may have multiple equilibria, limit cycles, or unbounded growth.

P4 — Mobility Capacity Bound
No other principle imposes an upper bound on density or requires mobility collapse at $\varrho_{\max}$. Without P4, the architecture has no horizon and no finite capacity.

P5 — Participation Feedback Loop
No other principle introduces memory, temporal accumulation, or feedback. Without P5, the system reduces to a single-variable diffusion with no oscillation or temporal causality.

P6 — Damping Discriminant
No other principle defines the algebraic boundary between oscillatory and monotonic regimes. Without P6, the architecture has no boundary motif and no regime partition.

P7 — Nonlinear Triad Coupling
No other principle generates higher-harmonic content or the invariant triad ratio. Without P7, the architecture loses its spectral fingerprint.

Conclusion:
Each principle introduces a structural feature that cannot be reconstructed from the others.
The Canon is independent.

4.2 Minimality Proof
To demonstrate minimality, we show that removing any one principle breaks at least one essential architectural layer — a motif, a law, a geometric object, or an invariant — and that no smaller set of principles can generate the full ED ontology.

Removing P1 (Operator Structure)
Breaks: Saddle, Dissipator, Triad, Modal Hierarchy, spatial causality, spectral causality, gradient sign invariant.
Result: The architecture collapses to uncoupled ODEs.

Removing P2 (Channel Complementarity)
Breaks: Boundary motif, regime structure, channel complementarity law, parameter-space geometry.
Result: No principled combination of direct and mediated influence.

Removing P3 (Penalty Equilibrium)
Breaks: Unique Attraction, Manifold, One-Dimensional Relaxation, direction of becoming, information erasure.
Result: No teleology; dynamics lose their target.

Removing P4 (Mobility Capacity Bound)
Breaks: Horizon motif, finite capacity law, bounded state space, kinetic saturation.
Result: Density becomes unbounded; no horizon geometry.

Removing P5 (Participation Feedback Loop)
Breaks: Spiral, Memory Loop, temporal causality, storage, delay, oscillatory regime.
Result: System collapses to monotonic diffusion.

Removing P6 (Damping Discriminant)

Breaks: Boundary motif, regime partition, critical damping law, spiral/monotonic distinction.
Result: Parameter space loses its topological structure.

Removing P7 (Nonlinear Triad Coupling)
Breaks: Triad motif, harmonic protection law, spectral fingerprint, invariant k_3/k_1 ratio.
Result: Architecture loses its nonlinear spectral identity.

Conclusion:
Removing any principle destroys essential architectural structure.
The Canon is minimal.

4.3 Counterfactual Analysis

To further demonstrate necessity, we examine the architecture that would result if each principle were removed while all others remained intact. These counterfactual architectures illustrate the fragility of the ED structure and the indispensability of each principle.

Without P1 — No Spatial Architecture
The system becomes a set of uncoupled temporal ODEs.
No geometry, no motifs involving gradients, no spectral structure.

Without P2 — No Regime Structure
The system has no boundary, no oscillatory/monotonic distinction, and no dual-channel dynamics.

Without P3 — No Teleology
The system may wander indefinitely, oscillate, or diverge.
No manifold, no attractor, no information erasure.

Without P4 — No Horizon
Density can grow without bound.
No finite capacity, no kinetic barrier, no saturation symmetry.

Without P5 — No Memory
The system loses oscillation, storage, delay, and feedback.
No spiral, no temporal causality.

Without P6 — No Boundary
The parameter space becomes topologically featureless.
No critical damping, no regime partition.

Without P7 — No Spectral Identity
The system becomes spectrally trivial.
No triad, no harmonic protection, no invariant ratio.

Conclusion:
Each principle is individually necessary.
The Canon is fragile in the precise sense required of a minimal axiom set.

5. Canonical Equivalence Class

The Architectural Canon does more than specify the structure of the unified ED cosmology PDE. It defines an entire equivalence class of dynamical systems — all PDEs that share the same seven canonical principles, regardless of their specific functional forms. This section formalizes that class, identifies what may vary without altering architectural identity, and clarifies what cannot vary without breaking the Canon.

5.1 Definition of the Canonical Equivalence Class

A PDE belongs to the ED canonical equivalence class if and only if it satisfies all seven canonical principles:

1. Operator Structure
2. Channel Complementarity
3. Penalty Equilibrium
4. Mobility Capacity Bound
5. Participation Feedback Loop
6. Damping Discriminant
7. Nonlinear Triad Coupling

Any PDE satisfying these principles — even with different mobility functions, penalty shapes, damping parameters, or domain geometries — will exhibit:

- the same eight motifs
- the same nine laws
- the same geometric objects
- the same causal channel
- the same information pipeline
- the same symmetry structure
- the same invariants

This is the architectural meaning of equivalence: different equations, same architecture.

Formally:

Two PDEs are architecturally equivalent if they satisfy the same canonical principles.

This equivalence is structural, not algebraic.

5.2 What Can Vary Without Breaking the Canon

The Canon is intentionally permissive about many details. A wide range of PDEs can instantiate the architecture as long as they preserve the seven principles. The following elements may vary freely:

Functional Forms

- The specific shape of the mobility function $M(\varrho)$
- The specific form of the penalty function $P(\varrho)$
- The exact dependence of M' on ϱ
- The smoothness or curvature of the nonlinear term

As long as:

- $M(\varrho_{max}) = 0$
- $P(\varrho^{\backslash *}) = 0$ with $P' > 0$
- $M'|\nabla\varrho|^2$ generates triad coupling

…the architecture remains intact.

Parameter Values

- Any $D \in [0,1]$
- Any $\zeta \geq 0$
- Any $\tau > 0$
- Any $\rho_{max} > \rho^{\backslash *}$

These parameters determine where the system sits in the architectural space, not what the architecture is.

Domain Geometry

- 1D, 2D, or 3D spatial domains
- Periodic, reflective, or open boundaries
- Arbitrary domain shapes

The Canon is agnostic to geometry; it governs the architecture of the dynamics, not the topology of the container.

Scaling and Multiplicative Factors

- Multiplying the entire operator by a positive scalar
- Rescaling time
- Rescaling the mobility or penalty functions

These transformations preserve the Canon because they preserve the structural relationships.

Conclusion:
The Canon defines a family of PDEs, not a single equation.

5.3 What Cannot Vary Without Breaking the Canon
The following elements are architecturally essential. Changing any of them destroys at least one canonical principle and therefore breaks the ED identity.

Operator Structure
Removing or altering the three-term structure of $F[\varrho]$ breaks P1.

Dual-Channel Dynamics
Eliminating the D/H split or violating $D+H = 1$ breaks P2.

Unique Equilibrium
Allowing multiple equilibria or non-monotone penalties breaks P3.

Finite Capacity
Removing the mobility collapse at $\varrho_{\max}$ breaks P4.

Memory Loop
Removing the participation variable or its feedback breaks P5.

Damping Discriminant
Altering the discriminant or removing the threshold breaks P6.

Triad Coupling
Removing the nonlinear gradient-squared term or altering its spectral consequences breaks P7.

Conclusion:
These elements are the identity conditions of the ED architecture.

5.4 Multiplicative Equivalence
A particularly important subclass of equivalence arises from multiplicative transformations:

- Multiplying the operator $F[\varrho]$ by a positive scalar
- Multiplying M, M', or P by a positive scalar
- Rescaling time or the participation variable
- Applying uniform amplitude scaling within the capacity bound

These transformations preserve:

- the operator structure
- the channel complementarity
- the equilibrium location
- the damping discriminant
- the triad ratio
- all invariants

Thus:

Any two PDEs related by multiplicative transformations are architecturally identical.

This explains why the ED architecture is robust to parameter variation and functional-form perturbation: the Canon defines a structural identity, not a numerical one.

5.5 Architectural Identity
A PDE is architecturally ED if and only if:

- it satisfies all seven canonical principles
- it generates the same motifs, laws, geometry, and invariants

- it belongs to the canonical equivalence class

This is the architectural meaning of identity:

> **The Canon defines what ED is.**
> **The equivalence class defines which PDEs count as ED.**

6. Canonical Form of the Architecture

The Architectural Canon defines a class of PDEs, not a single equation. Yet every PDE in this class can be expressed in a canonical form that makes the seven principles explicit. This section presents that form: the canonical operator, channels, equilibrium, discriminant, coupling, and state-space structure. Together, these elements constitute the architectural "normal form" of the ED universe.

6.1 Canonical Operator

At the heart of the architecture is the canonical operator:

$$F[\rho] = M(\rho)\nabla^2\rho + M'(\rho)|\nabla\rho|^2 - P(\rho)$$

This operator expresses three irreducible structural commitments:

1. Curvature-dependent diffusion (via $M\nabla^2\varrho$)
2. Nonlinear gradient coupling (via $M'|\nabla\varrho|^2$)
3. Monotone restoring force (via $-P(\varrho)$)

These terms correspond directly to Canonical Principle 1 and encode the architecture's spatial, spectral, and restorative structure.

6.2 Canonical Channels

Density evolves through two complementary channels:

$$\dot{\rho} = D\,F[\rho] + H\,v, \qquad D + H = 1, \qquad D,H \in [0,1].$$

This dual-channel structure expresses Canonical Principle 2:

- The direct channel $D\,F[\varrho]$ applies the operator immediately.
- The mediated channel $H\,v$ applies it through memory.

The conservation constraint $D+H = 1$ defines the architecture's parameter space and ensures that all dynamics lie on a one-simplex.

6.3 Canonical Equilibrium

The restoring force $P(\varrho)$ has a single zero:

$$P(\rho^*) = 0, \qquad P'(\rho) > 0.$$

This expresses Canonical Principle 3:

- The equilibrium is unique.
- It is globally attracting.
- It defines the direction of becoming.

The canonical equilibrium point in phase space is:

$$(\rho,v) = (\rho^{*},0).$$

All trajectories converge to this point through either spiral or monotonic relaxation.

6.4 Canonical Discriminant
The qualitative flow type is determined by the canonical discriminant:

$$\Delta = D+2\zeta.$$

This expresses Canonical Principle 6:

- Under-damped when $\Delta < 1$
- Critically damped when $\Delta = 1$
- Over-damped when $\Delta > 1$

The discriminant defines the Boundary Surface in parameter space:

$$D + 2\zeta = 1.$$

This boundary partitions the architecture into two topologically distinct dynamical regimes.

6.5 Canonical Coupling
The nonlinear triad coupling is encoded in the gradient-squared term:

$$M'(\rho)\,|\nabla\rho|^2.$$

This expresses Canonical Principle 7:

- It generates higher harmonics, especially $k = 3$ from $k = 1$.
- It produces the invariant triad ratio k_3/k_1.
- It defines the architecture's spectral fingerprint.

The triad coupling is the source of the Nonlinear Triad motif, the Harmonic Protection law, and the Triad Locus geometry.

6.6 Canonical State Space
The mobility capacity bound defines the architecture's state-space geometry:

$$M(\rho_{max}) = 0, \qquad M(\rho) > 0 \quad \text{for } \rho < \rho_{max}.$$

This expresses Canonical Principle 4:

- Density is bounded above.
- Diffusive transport collapses near ϱmax.
- A kinetic horizon forms where mobility becomes small.

Thus the canonical state space is:

$$0 \le \rho \le \rho_{max}, \qquad v \in R.$$

The horizon at ϱ_{max} is a structural feature of the architecture, not a boundary condition.

6.7 Canonical System Summary
The canonical form of the ED architecture can be summarized as the coupled system:

$$\begin{aligned} \dot{\rho} &= DF[\rho] + Hv, \quad D + H = 1, \\ \dot{v} &= \frac{1}{\tau}(F[\rho] - \zeta v), \\ F[\rho] &= M(\rho)\nabla^2\rho + M'(\rho)|\nabla\rho|^2 - P(\rho), \\ P(\rho^*) &= 0, \quad P'(\rho) > 0, \quad M(\rho_{\max}) = 0. \end{aligned}$$

This system is the canonical representative of the entire architectural equivalence class defined in Section 5.

7. Unified Interpretation

The Architectural Canon does more than specify the minimal axioms of the ED architecture. It reveals the deep structural logic that makes the architecture coherent, universal, and inevitable. By examining the Canon as a whole, we can understand not only how the ED cosmology PDE behaves, but why it behaves that way — and why any system satisfying the same seven principles must exhibit the same motifs, laws, geometry, and invariants.

7.1 What the Canon Reveals About ED

(a) ED is defined by structure, not by equations
The Canon shows that ED is not a specific PDE, nor a particular choice of mobility or penalty functions. It is a structural architecture defined by seven irreducible commitments. Any PDE that satisfies these commitments is architecturally ED, regardless of its algebraic details.

This elevates ED from a model to a theory class.

(b) The architecture is generative
The Canon explains why:

- the eight motifs appear and no others
- the nine laws hold universally
- the geometric objects have their specific shapes
- the causal channels have their specific roles
- the information pipeline has its nine stages

- the invariants are robust under perturbation

These features are not empirical observations — they are logical consequences of the Canon.

(c) The architecture is internally coherent
The Canon reveals the deep interdependence of the architectural layers:

- P1 and P7 generate the spectral structure
- P2 and P5 generate the temporal structure
- P3 and P4 generate the state-space structure
- P6 organizes the parameter-space structure

The architecture is not a collection of unrelated features; it is a tightly coupled system in which each principle reinforces the others.

(d) The architecture is inevitable
Given the seven principles, the entire ED ontology follows with no additional assumptions.
This inevitability is the hallmark of a well-formed architectural theory.

7.2 Why the Canon Explains Universality
The Canon clarifies why ED dynamics are universal across:

- initial conditions
- parameter choices
- functional-form perturbations
- multiplicative transformations
- domain geometries

The reason is structural:

> **Universality arises because the Canon defines the architecture, not the equation.**

As long as the seven principles hold:

- the same motifs appear
- the same laws govern the dynamics
- the same invariants remain fixed
- the same geometric objects shape the flow
- the same causal and informational structures operate

This is architectural universality — the strongest form of universality in dynamical systems.

7.3 Why the Canon Explains Fragility
The Canon also explains why ED is fragile in a precise, structural sense:

- Removing P1 destroys spatial dynamics
- Removing P2 destroys regime structure

- Removing P3 destroys teleology
- Removing P4 destroys finite capacity
- Removing P5 destroys memory and oscillation
- Removing P6 destroys the boundary
- Removing P7 destroys the spectral fingerprint

Each principle is individually necessary.
The architecture has no redundancy.

This fragility is not a weakness — it is the defining feature of a minimal axiom set.

7.4 Closure of the Foundational Arc
The ED-Arch series began with motifs (ED-Arch-02) and progressively extracted deeper structural layers until reaching invariants (ED-Arch-10) and synthesis (ED-Arch-11). The Canon completes this progression by identifying the axiomatic root from which all layers descend.

The foundational arc is now closed:

$$\text{Canon} \longrightarrow \text{Laws} \longrightarrow \text{Motifs} \longrightarrow \text{Geometry} \longrightarrow \text{Dynamics}.$$

The Canon is the apex of the ED-Arch hierarchy and the foundation of the ED cosmology PDE.

7.5 Implications for Future Work
The Canon provides the conceptual foundation for two major research directions:

(a) ED-Arch-X: Applications of the Architecture
Using the Canon to classify, diagnose, and design other dynamical systems.

(b) ED-Cosmo: Physical Interpretation
Using the Canon to interpret cosmological phenomena — expansion, structure formation, horizons, dimensionality — through the ED architecture.

Both arcs rely on the Canon as their structural anchor.

7.6 Summary
The Architectural Canon reveals that:

- ED is defined by seven irreducible principles
- These principles generate the entire architecture
- The architecture is coherent, universal, and inevitable
- The Canon defines the identity conditions for ED universes
- The foundational arc of ED-Arch is complete

The Canon is not merely a summary — it is the architectural essence of the ED cosmology PDE.

8. Conclusion

The Architectural Canon identifies the seven irreducible principles that define the Event-Density architecture at its deepest structural level. These principles are not empirical observations, nor are they summaries of the ED-Arch series; they are the axioms from which the entire architecture follows. Every motif, law, geometric object, causal channel, information transformation, symmetry, and invariant described in ED-Arch-02 through ED-Arch-11 is a logical consequence of these seven commitments.

The Canon demonstrates three essential facts about the ED cosmology PDE:

1. **Identity:**
    - ED is not a particular equation but a structural architecture.
    - Any PDE satisfying the seven canonical principles is architecturally ED.
2. **Universality:**
    - The architecture is robust to changes in functional forms, parameter values, and domain geometries.
    - The Canon explains why diverse systems exhibit the same motifs, laws, and invariants.
3. **Fragility:**
    - Each principle is individually necessary.
    - Removing any one breaks specific architectural layers and destroys the ED identity.

With the Canon established, the foundational arc of the ED-Arch series is complete. The architecture now stands on a minimal, irreducible basis, and the unified ED cosmology PDE can be understood as a specific instantiation of this deeper structure. The Canon provides the conceptual foundation for all future work — whether extending the architecture to new systems (ED-Arch-X), interpreting physical cosmology through ED (ED-Cosmo), or developing simulation frameworks that instantiate the architecture in practice.

The ED architecture is now fully specified:

seven principles, one identity, one universe of systems.